



# Introduction

Distributed Systems  
Bachelor In Informatics Engineering  
Universidad Carlos III de Madrid


# What we'll see in this class



- The evolution of computer science
- What is a distributed system
- Examples of distributed applications
- Strengths and weaknesses of distributed systems
- Distributed vs. parallel systems
- Design challenges
- What is the *middleware*
- Computation paradigms for distributed systems
- Hardware platforms for distributed systems

# The evolution of computer science

- The **70s**:
  - Mainframes
    - Shared time systems
    - Centralized resources
    - Simple terminals
  - User interfaces were not very friendly.
  - First computer networks.



# The evolution of computer science

- The **80s**:
  - PCs and workstations
- Applications are mostly complex and executed locally
- UIs become more friendly
- Local networks (LAN)
- First distributed Oss
  - Mach, Sprite, Chorus, ...




# The evolution of computer science

- The 90s:
  - Client/server applications
  - Execution becomes more decentralized
  - The web!
  - Web-based applications take off
    - E-commerce
    - Multimedia
    - Control systems
    - Medical applications
    - Supercomputing using the Internet

# The evolution of computer science


- **2000:**

- New distributed computing paradigms
  - Grid computing
  - Peer-to-Peer
  - Ubiquitous computing
- Mobile devices
- Web-based applications
- Mostly network-based applications




# The evolution of computer science

- Traditional infrastructure




# The evolution of computer science

- Traditional infrastructure



- Resource consolidation






# Evolution


- **Bell's Law (1972):**
  - A new computer technology will come out every 10 years




- **Moore's Law (1970):**
  - The number of transistors per chip will double every 24 months



# Distributed Systems



- Consists of physically distributed resources (hardware and software) which are connected via a network and are able to collaborate to accomplish a given task.



# Recap



- A process is an executing program
- A computer network is a set of computers connected via an interconnection network
- A distributed system is a set of computers (w/o shared memory nor common clock) which are connected via an interconnection network
  - Distributed application: A set of processes executing on one or more computers and which collaborate and communicate via message passing
- A protocol is a set of rules and instructions which specify how communication is realized in a distributed system via message passing

# Another definition of a distributed system

*"A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable."*

-Leslie Lamport

# Examples: The Web

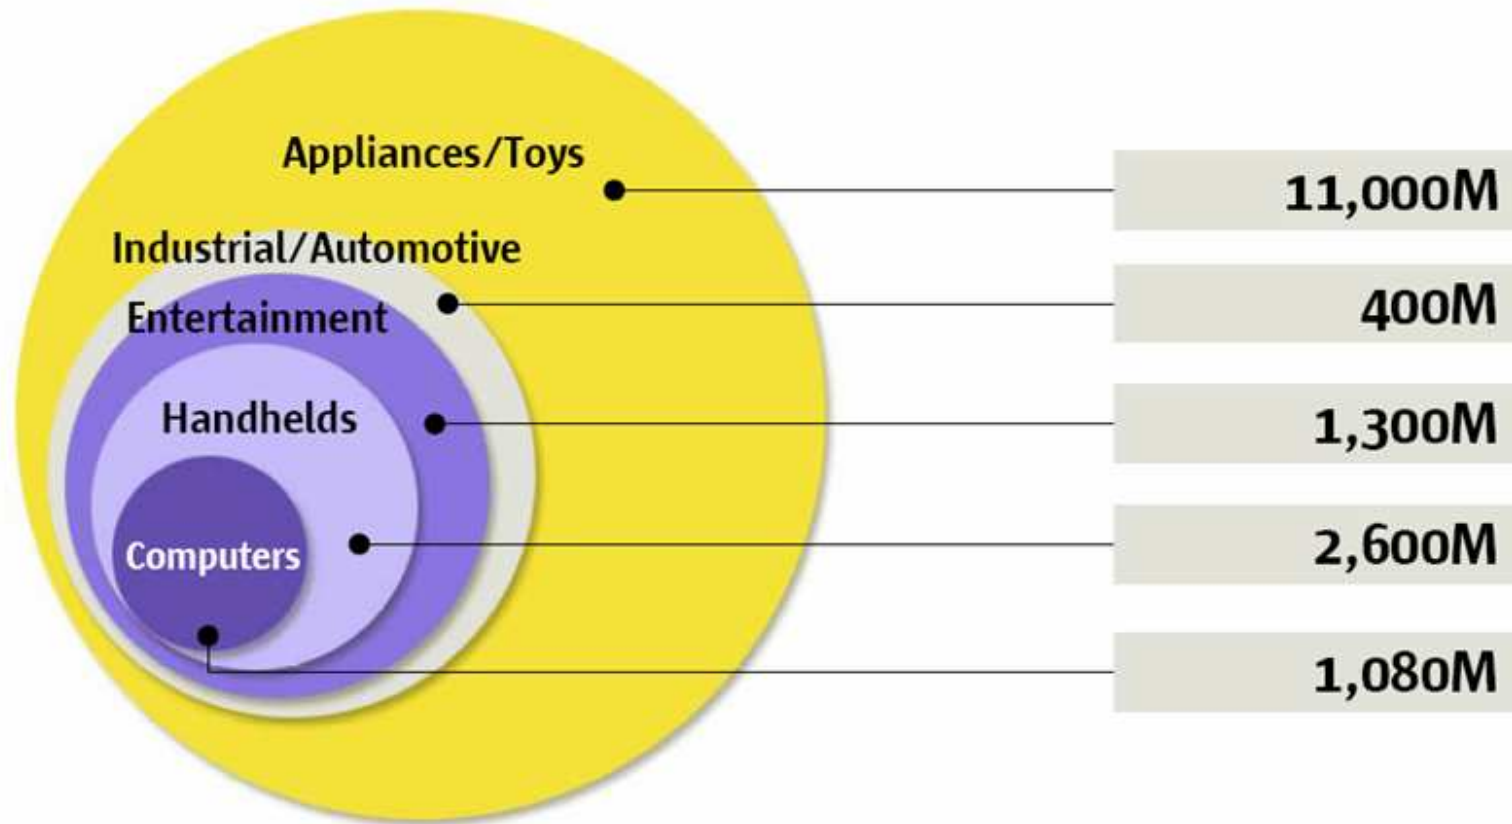


# Examples of distributed systems and applications

- ▶ Email (IMAP, POP)
- ▶ File transfer programs (FTP)
- ▶ News services
- ▶ World Wide Web (WWW)
- ▶ Air traffic control systems
- ▶ Bank applications
- ▶ E-commerce
- ▶ Multimedia applications (videoconferences, video on-demand , etc.)
  - ▶ Bandwidth is an order of magnitude bigger than for other apps.
  - ▶ They require quality of service (QoS)
- ▶ Medical apps. (image transfer)

# In the future...

WW Installed, 2012



Source: IDC Estimates, 2004



# Advantages of distributed systems

- ▶ Resource sharing (HW, SW, data)
  - ▶ Remote resource access
    - ▶ Client-server model
    - ▶ Object-based model
- ▶ Good cost/performance ratio
- ▶ Scalable
- ▶ Fault-tolerance
  - ▶ Replication
- ▶ Concurrency: multiple simultaneous users
- ▶ Computing speed/capacity for executing apps. on many machines in parallel



# Weaknesses of distributed systems

- Interconnection:
  - Cost
  - Reliability, message loss
  - Saturation
- Security of communication may be a problem
- Software is more complex

# Distributed vs. concurrent systems


- **Distributed systems**

- Goal: resource sharing, collaboration
- Computers connected in a network


- **Concurrent systems**

- Goal:
  - High performance(*speedup*)
  - High productivity
- Parallel machines (dedicated architectures)
  - Multiprocessors, clusters, supercomputers
- Networks of workstations working in cluster mode
- Grid Computing ([www.gridcomputing.com](http://www.gridcomputing.com))


# Parallel architectures



Shared memory




Distributed shared memory




Distributed memory

# Network topologies






Grid



Tree



Linear



# Distributed systems – design challenges

- Components are heterogenous
- Naming issues
- Communication and synchronization issues
- Performance issues
- Concurrency
- Scalability
- Software structure
- Reliability
- Quality of service (QoS)
- Transparency

# Heterogeneity

- Components in a distributed system may vary in terms of:
  - Networks
  - HW
  - OS
  - Programming languages
  - Applications

# Solution

- ▶ Open systems
  - ▶ Public specs and interfaces (e.g. RFCs)
  - ▶ Uniform communication mechanisms
  - ▶ Can be built on top of heterogenous HW and SW
- ▶ Examples:
  - ▶ TCP/IP
  - ▶ NFS
  - ▶ CORBA ([www.omg.org](http://www.omg.org))
  - ▶ Globus ([www.globus.org](http://www.globus.org))
  - ▶ Web services

# Naming conventions

- Users **name** objects (e.g. [www.uc3m.es](http://www.uc3m.es))
- Programs name objects via an **identifier** (e.g. 163.117.131.31)
- **Resolving a name** means obtaining the identifier starting from the name
- Important: the names must be domain-independent
- Design issues to consider:
  - The namespace (size, structure, hierarchy, ...)
  - The naming system which does the name resolution (e.g. DNS)



# Communication and synchronization

- ▶ Basic mechanism: message passing
  - ▶ Synchronous
  - ▶ Asynchronous
- ▶ Processes on different machines communicate via send/receive primitives
  - ▶ Remote procedure calls
  - ▶ Remote object invocation
- ▶ Group communication
  - ▶ *Multicast, broadcast*
  - ▶ Useful for teamwork, fault tolerance, replication, ensuring data consistency, etc

# What to take into consideration

- For a system with  $n$  users to be scalable the number of resources must be proportional to  $O(n)$
- Use distributed algorithms
  - Generally these make use of hierarchical rather than linear structures
- Avoid bottlenecks
  - Decentralized algorithms
- Avoid the SW resources running out
  - E.g.: 32 bits for IP addresses

# Software structure

Centralized systems



- ▶ The OS:
  - ▶ manages the HW resources efficiently
  - ▶ offers services to the applications


# Software structure

Distributed systems

- Three options:
  - Use machines with separate OSs
  - Use a distributed OS
  - Use a middleware or a distributed environment
- Important to make programming distributed apps easy and transparent

# Machines with separate OSs

- The user sees a set of independent machines
  - No transparency
- Access to other machine's resources is explicit
- Hard to use



# Reliability

- The probability that a system keeps delivering the desired functionality under given conditions and during a given period of time
- To ensure reliability one must guarantee:
  - Consistency
  - Security
  - Error handling

# Consistency

- The problem occurs when various processes access and update data concurrently – aka coherence
  - Update consistency
  - Replication consistency
  - Cache coherence
  - Failure consistency
  - Clock consistency

# Failure handling

- In general in distributed systems partial failures may occur
- The goal of a distributed system is availability
  - Measures the percentage of time that a system is available for use
- Improving availability:
  - Tolerating failures
  - Failure detection
  - Failure masking
  - Failure recovery
  - Redundancy



# Quality of service (QoS)

- ▶ The ability to satisfy the time requirements when transmitting and processing multimedia data streams in real time
- ▶ System performance is measured in:
  - ▶ Response time
    - ▶ Latency
  - ▶ Data transfer rate
    - ▶ The transfer speed between two computers, usually measured in bps
- ▶ The performance is determined by:
  - ▶ Communication network
  - ▶ Communication services
  - ▶ OS
  - ▶ The distributed programming support

# Transparency

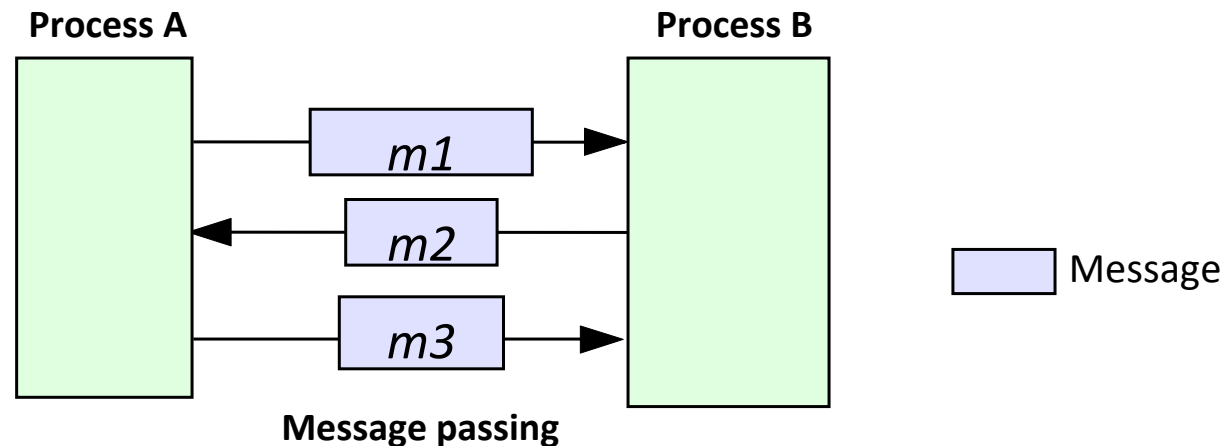
- The concealment from the user of the separation of components in a distributed system
  - **Access T**: local and remote resources are accessed via same operations
  - **Location T**: access without knowledge of their physical/network location
  - **Concurrency T**: processes may operate concurrently on shared resources without interference
  - **Replication T**: multiple instances of resources to increase reliability and performance
  - **Failure T**: concealment of faults
  - **Mobility T**: allows movement of resources and clients without affecting operation of users and programs
  - **Scaling T**: allows system and apps to expand without changes in system structure and algorithms

# Distributed computation paradigms

- Message passing
- Client-server
- Remote procedure call
- Peer-to-peer
- Distributed objects
- Mobile agents
- Services
- Collaborative applications (*groupware*)


# Message passing

- The fundamental paradigm for building distributed apps
- A process sends a request message
- The message is received by the destination process which processes it and replies with another message
- The reply may cause further messages from the sender process



# Client-Server

- Assigns different roles - client and server – to the communicating processes
- Server:
  - Offers a service
  - Passive: waits for requests
  - May in turn be clients of other services
- Client:
  - Asks for service
  - Active: sends requests

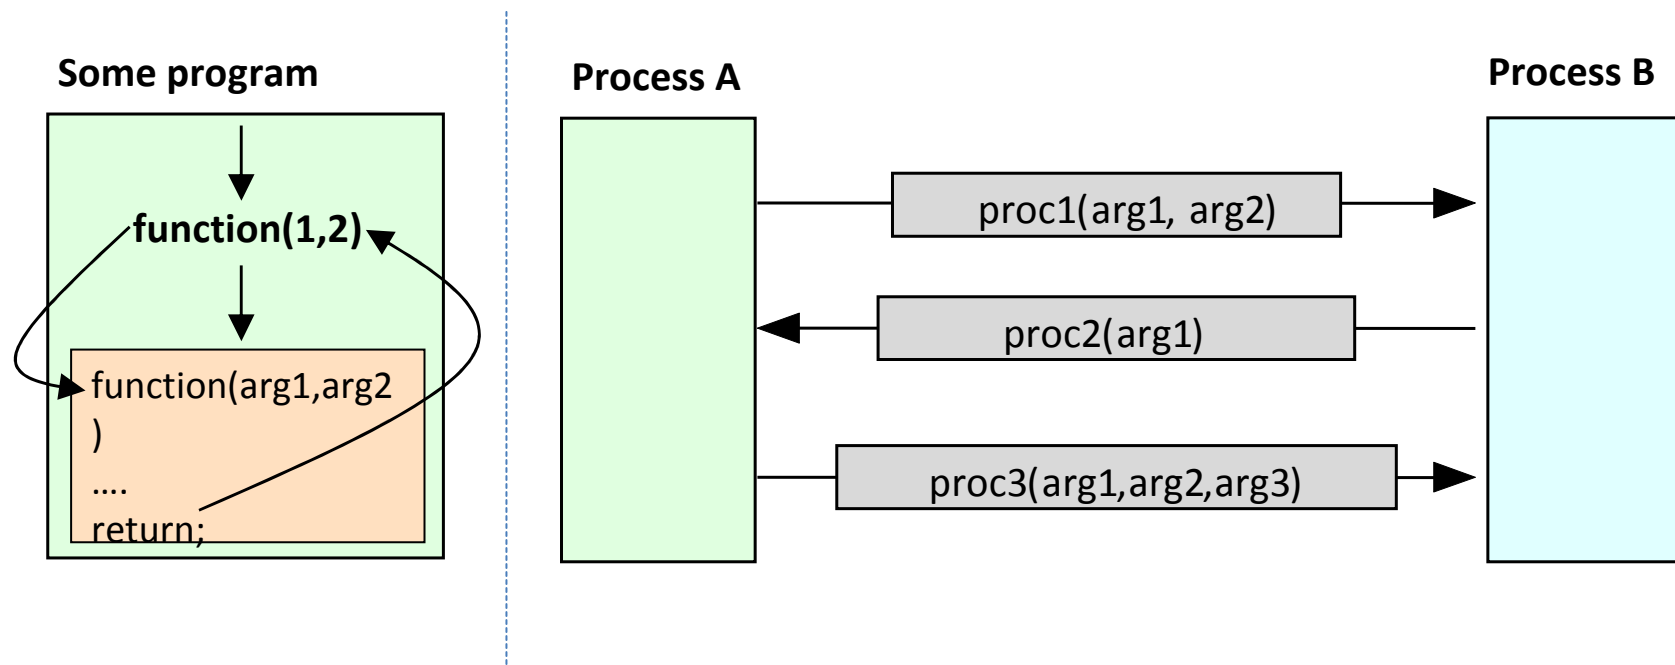


# Client-Server

- Efficient abstraction
- Asymmetrical role assignment helps with synchronization
- Implemented via:
  - Sockets
  - Remote procedure calls (RPC)
  - Remote method invocation (RMI, CORBA, ...).
- Mostly useful for centralized services
- E.g.: Internet services (HTTP, FTP, DNS, ... )


# Remote procedure calls

- Idea: make it such that distributed apps are programed the same as non-distributed ones
- Conceptually same as a local procedure call



# Remote method invocation


- Similar to RPCs – the process invokes a method local to another process
- **E.g.:** CORBA, Java RMI, Microsoft COM, DCOM, Java Beans, .NET Remoting






# Peer-to-Peer

- ▶ The processes involved in communication have the same role:
  - ▶ Client **and** server
  - ▶ Efficient access to data and other resources that they collectively store and manage
- ▶ E.g.:
  - ▶ *Napster* → file sharing




# Mobile agents

- ▶ A running program that
  - ▶ Is launched on an origin computer
  - ▶ Travels from one computer to another in a network
  - ▶ Carries out a task on someone's behalf
  - ▶ Eventually returns with the results
- ▶ No message exchange
- ▶ Security problem: the executable mobile code may be malicious




# Web services

- A process requests a service via a reference provided by the directory service
- Services must register with the directory service to get their services published
- Location transparency
- E.g.: **SOAP**




# Collaborative applications (*groupware*)

- Several processes participate in a common work session
- Communication: unicast, multicast, broadcast
- Two types:
  - Message-based: send messages to others in the group
  - Whiteboard-based: use a shared whiteboard to R/W data



***Message-based groupware***



***Whiteboard-based groupware***