

Synchronization and communication



ARCOS Group

Distributed Systems

Bachelor In Informatics Engineering


Universidad Carlos III de Madrid

Syllabus

- ▶ Processes and threads
- ▶ Concurrency
- ▶ Synchronization mechanisms
- ▶ Communication mechanisms
- ▶ POSIX

Distributed system

- ▶ Consists of physically distributed resources (hardware and software) which are connected via a network and are able to collaborate to accomplish a given task



Monothread vs. Multithread processes

- ▶ A lightweight process / thread is an execution flow which shares the memory image (and other information) with other lightweight processes
 - ▶ Depending on the number of processes that can execute simultaneously an OS is unitask or multitask
- ▶ A process may consist of a single execution flow (aka monothread) or more (aka multithread)


Private and shared information

▶ Per *thread* (private)

- ▶ PC, registers
- ▶ Stack
- ▶ Status (executing, ready, blocked)


▶ Per process (shared)

- ▶ Space in memory
- ▶ Global variables
- ▶ Open files
- ▶ Child processes
- ▶ Timers
- ▶ Signals and semaphors
- ▶ Statistics



The status of a lightweight process

- ▶ Just like a conventional process, a lightweight process may be in the following states: executing, ready, or blocked (or suspended for medium-term scheduling)
 - ▶ The status of a process is the combination between the states of the lightweight processes that it consists of



Threads and concurrency

- ▶ Lightweight processes enable easy app. paralelization
 - ▶ A program may be split into parts which may execute concurrently within lightweight threads
 - ▶ Encourages modular application design
- ▶ The basic idea is to always keep some lightweight process executing – while a thread is blocked another one may execute


```
main() {  
  ...  
  function1(args);  
  function2(args);  
  ...  
}
```

Sequential design




```
main() {  
  ...  
  run_thread (function1(args));  
  run_thread (function2(args));  
}
```

Parallel design




E.g.: concurrent servers with *threads*

- ▶ The concurrent processes must communicate and synchronize



Concurrent design



Sequential design


Server design using *threads*

Several alternatives to build parallel servers:

A single process which accepts requests and either (1) distributes them to threads from a thread pool or (2) creates a new thread to service the request

Set of similar threads which can read requests from a port

Pipeline the work and have a specialized thread for each stage



Programming using threads

- ▶ UNIX offers the following thread management operations:
 - ▶ Create and identify threads
 - ▶ Associate attributes with threads
 - ▶ Schedule threads
 - Priorities, scheduling policies
 - ▶ Terminate threads
 - Does NOT imply the process' children must terminate immediately
 - Parent process may wait for children to finish – wait
 - Process may wait for another process to finish - join

Create and identify *threads*

► Creation:

```
int pthread_create(pthread_t* thread,  
    const pthread_attr_t* attr, void * (*function)(void *),  
    void * arg);
```

where:

thread	new thread id
attr	attributes of the new thread (or NULL)
function	the function that the new thread will execute
arg	pointer to the thread arguments

returns:

int	an integer if successful creation; -1 otherwise
-----	---

► Getting the thread id:

```
pthread_t pthread_self(void)
```

Thread attributes

- ▶ Default values

- ▶ **E.g.:** a thread is *joinable* (not independent) if not defined otherwise

- ▶ To modify the thread attributes:

- ▶ Initialize an attribute object to use it when creating new threads

```
int pthread_attr_init(pthread_attr_t *attr);
```

- ▶ Destruct an attribute

```
int pthread_attr_destroy(pthread_attr_t *attr);
```

- ▶ Change termination status (Joinable / Detached)

```
int pthread_attr_setdetachstate(  
                                pthread_attr_t *a, int detachstate);
```

If detachstate is

- **PTHREAD_CREATE_DETACHED** the process is independent
 - **PTHREAD_CREATE_JOINABLE** the process is NOT independent

- ▶ Other attributes: stack size, scheduling policy, priority, etc

Terminating a thread

- ▶ Finalize the thread execution

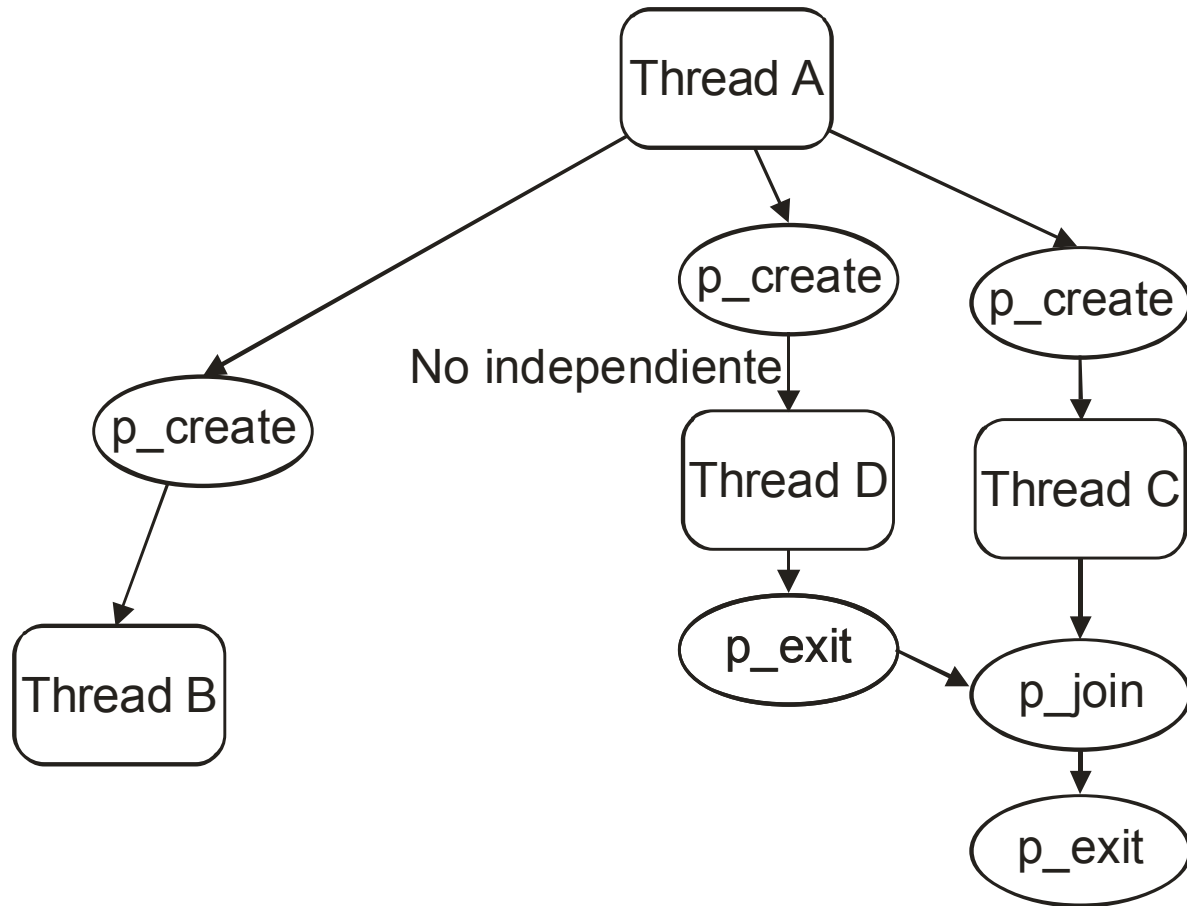
```
int pthread_exit(void *value);
```

- ▶ If a process is **not independent (joinable)** then need to call

```
int pthread_join(pthread_t thid, void **value);
```

- ▶ Wait for process with id *thid* to finish
- ▶ Return in second argument the value passed in *pthread_exit*.
- ▶ Can't wait for termination of independent threads

E.g.: Thread hierarchy



E.g. I: Creation and termination

```
#include <stdio.h>    /* printf */
#include <pthread.h> /* For threads */
#define NUM_THREADS 10

int function(int *idThread){
    printf("Thread id = %d\ti=%d: ", pthread_self(),*idThread);
    pthread_exit(NULL);
}

int main () {
    pthread_t arrayThread[NUM_THREADS]; /* Array of threads */
    int i;
    /* CREATE THREADS */
    for(i=0;i<NUM_THREADS;i++)
        if (pthread_create(&arrayThread[i], NULL, (void *)funcion, &i)==-1)
            printf("Error\n");
    /* WAIT FOR TERMINATION */
    for(i=0;i<NUM_THREADS;i++)
        pthread_join (arrayThread[i], NULL);
    exit(0);
}
```

E.g. II: Modifying attributes


```
#include <stdio.h>    /* printf */
#include <pthread.h> /* For threads... */
#define MAX_THREADS 10

void func(void) {
    printf("Thread %d \n", pthread_self());
    pthread_exit(0);
}

main() {
    int j;
    pthread_attr_t attr;
    pthread_t thid[MAX_THREADS];

    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);
    for(j = 0; j < MAX_THREADS; j++)
        pthread_create(&thid[j], &attr, func, NULL);
    sleep(5);
    pthread_attr_destroy(&attr);
}
```


Concurrency





► Models:

- Multiprogramming using one processor
- Multiprocessor (shared memory system)
- Multicomputer (distributed system)

► Advantages:

- Faster processing
- Multiuser systems w interaction
- Better use of resources

Monoprocessor vs. Multiprocessor



Types of concurrent processes


- ▶ **Independent:** successful execution does not depend on collaborating with other processes
- ▶ **Cooperating:** collaboration is necessary to perform their function
- ▶ In both cases you have process interaction
 - ▶ they compete for resources
 - ▶ they share resources

Recap




- ▶ A **race condition** involves unsafe (uncoordinated) concurrent access to shared resources
- ▶ A **critical section (CS)** is a code fragment via which various processes have access to shared resources
 - ▶ They may only execute on the processor on which they are entered
 - ▶ Usually enforced via semaphores
- ▶ An **atomic section** is an indivisible and uninterruptible code fragment
 - ▶ Guarantees that the invariants are observed and preserved by all operations
 - ▶ Often has a succeed-or-fail semantics
 - ▶ Usually enforced via mutual exclusion; may also be enforced via lock-free mechanisms, e.g. transactional memory

E.g.: Concurrent server



E.g.: Accessing a shared cache




Race condition: `Update_block()`

```
Update_block(block) {  
    b = Get_block();  
    copy(b, block);  
}
```

```
copy(b1, b2) {  
    for (i=0; i<length;i++)  
        *b1++ = *b2++;  
}
```

Accessing a shared cache



Race condition: `Update_block()`

```
Update_block(block) {  
    b = Get_block();  
    copy(b, block);  
}
```

```
copy(b1, b2) {  
    for (i=0; i<length;i++)  
        *b1++ = *b2++;  
}
```

Non atomic operation

Critical section

- ▶ n processes
 - ▶ Each one of them has a code fragment: **critical section**
 - ▶ Can't have more than 1 process at a time in a critical section
- ▶ To solve the critical section problem you must:

Enter CS

CS code

Exit CS

- ▶ Any mechanism solving the problem must ensure:
 - ▶ **Mutual exclusion**
 - ▶ **Progress**
 - ▶ **Bounded wait (no starvation)**

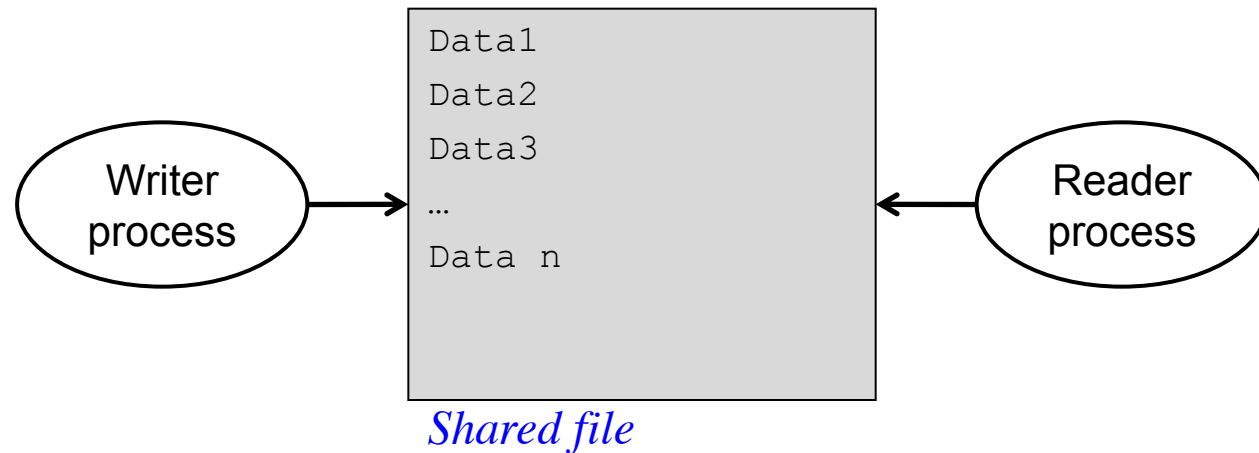
Communication mechanisms

- ▶ Files
- ▶ *Pipes* / FIFOs
- ▶ Shared memory variables
- ▶ Message passing




Files

- ▶ May be used for communicating between processes
 - ▶ Some write others read




Pipes

- ▶ A **pipe** is a communication and synchronization mechanism
 - ▶ Communication: shared file
 - ▶ Synchronization: due to the blocking semantics
- ▶ Aka a pseudo-archive identified by two descriptors: one for R, one for W
 - ▶ Same mechanism for R and W as for files
 - ▶ Atomic R and W



Synchronization mechanisms

- ▶ Signals (asynchronous communication)
- ▶ Pipes / FIFOs
- ▶ Semaphores
- ▶ Mutex and condition variables
- ▶ Message passing
- ▶ Monitors



Semaphores

- ▶ A **semaphore** [Dijkstra, 1965] is a synchronization mechanism for processes which execute on the same machine
- ▶ A protected variable which provides an abstraction for controlling access to a single resource
 - ▶ Keeps track of the number of resources available at the time
 - ▶ Provides operations to adjust the number of resources in a race-free manner (and wait if no resource is available)
- ▶ Processes are trusted to follow the protocol. If not problems may occur!
- ▶ Even if they do follow the protocol, multi-resource deadlock may occur if different semaphore / resource

Semaphores (cont'd)

- ▶ Two atomic operations:
 - ▶ `P` - aka `wait`
 - ▶ `V` - aka `signal`
- ▶ Atomicity achieved via atomic primitive HW instruction or mutual exclusion algorithm
 - ▶ On uniprocessors: disable interrupts is enough

Operations on semaphores

► P wait

```
wait(s)
{
    s = s - 1;
    if (s < 0) {
        <block the process>
    }
}
```

► V signal

```
signal(s)
{
    s = s + 1;
    if (s <= 0) {
        <unblock a process
        blocked on wait>
    }
}
```

s: integer value associated with the semaphore

Use of semaphores to protect the critical section :

wait(s)


Critical section

signal(s)

Critical sections with semaphores

```
wait(s); /* enter critical section */  
< critical section >  
signal(s); /* exit critical section */
```

Use a binary semaphore (with initial value 1)




Producer-consumer

Use sleep and wakeup routines

Sleep blocks the caller until another process wakes it up

```
Producer() {  
    while (1) {  
        Produce item  
        if buffer_full() sleep()  
        insert item  
        counter++  
        if (counter==1) wakeup(Consumer)  
    }  
}  
  
Consumer() {  
    while (1) {  
        if buffer_empty() sleep()  
        remove item  
        counter--  
        if (counter==size-1)  
            wakeup(Producer)  
        consume item  
    }  
}
```



Producer-consumer

- ▶ Race condition!
 - ▶ Consumer interrupted before calling sleep()
 - ▶ Producer wakes up consumer, BUT
 - ▶ Consumer not yet sleeping so wakeup call is lost
 - ▶ Consumer never awakened again (counter won't be 1 again!)

Producer-consumer with semaphores

```
Producer() {
    while (1) {
        Produce item
        wait(emptyCount)
        insert item
        signal(fillCount)
    }
}

Consumer() {
    while (1) {
        wait(fillCount)
        remove item
        signal(emptyCCount)
        consume item
    }
}
```

```
Producer() {
    while (1) {
        Produce item
        wait(emptyCount)
        lock(mutex)
        insert item
        unlock(mutex)
        signal(fillCount)
    }
}

Consumer() {
    while (1) {
        wait(fillCount)
        lock(mutex)
        remove item
        unlock(mutex)
        signal(emptyCCount)
        consume item
    }
}
```


- OK solution for 1 producer and 1 consumer
- For multiple: race condition between producers or consumers
- Solution: Execute critical section with mutual exclusion!!

Mutex

- ▶ A **mutex** is a binary semaphore
 - ▶ Used to describe a construct which prevents multiple processes from executing the same code (mutual exclusion in a critical section) / accessing same data at a given time
 - ▶ Usually has an owner – can't unlock if I'm not the one who locked the resource
- ▶ Implements two atomic operations
 - ▶ **lock(m)** – aka wait
 - ▶ If the mutex is already locked the process suspends
 - ▶ **unlock(m)** – aka signal
 - ▶ If there exists processes blocked waiting wake up one of them

Implementing critical sections using mutex

```
lock(m);    /* enter critical section*/  
< critical section >  
unlock(m);  /* exit critical section*/
```



Implementing critical sections using mutex

- ▶ Problem – not an efficient way to wait! (waste of CPU cycles)
- ▶ We need
 - ▶ A way for a thread to wait for something w/o wasting CPU cycles
 - ▶ A way for a thread to wake up another thread with some signal
- ▶ This abstraction is called a **condition variable**

Condition variables

- ▶ Synchronization variable associated with a mutex
- ▶ Implements two atomic operations:
 - ▶ **c_wait(cond,mutex)**
 - ▶ Blocks the calling thread and frees the mutex – aka thread goes to sleep
 - ▶ **c_signal(cond)**
 - ▶ Unblocks one or more sleeping processes waiting for the condition
 - ▶ The processes that wake up and the calling thread compete again for the mutex
- ▶ Blocking c.v. (Hoare-style) – signaled thread has priority
- ▶ Non-blocking c.v. – signaling thread has priority
 - ▶ Signal operation called notify

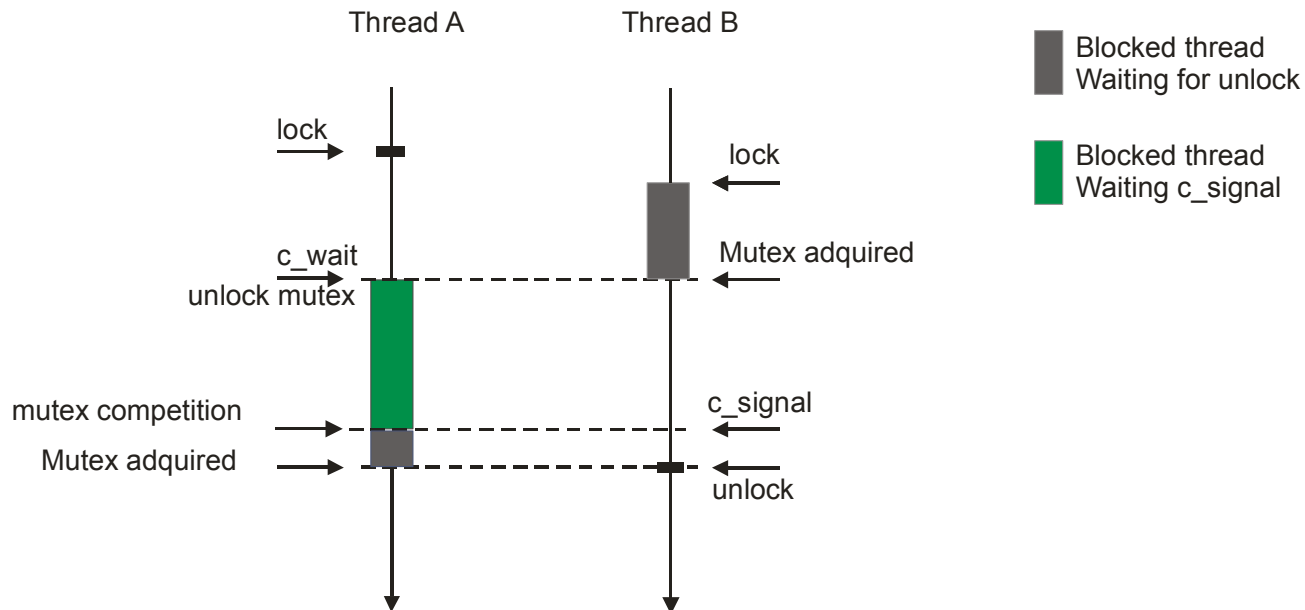
Using mutex and condition variables

Thread A

```
lock(m) ;  
/* Critical section */  
while (condition == FALSE)  
    c_wait(c,m) ;  
/* critical section */  
unlock(m) ; /* exit critical sect */
```

Thread B

```
lock(m) ;  
/* Critical section */  
/* modify condition */  
condition = TRUE;  
c_signal(c) ;  
/* critical section */  
unlock(m) ; /* exit critical section */
```



Using mutex and condition variables

▶ Thread A

```
lock(mutex); /* access resource */  
code;  
while (resource busy)  
    wait(condition,mutex);  
mark resource busy;  
unlock(mutex);
```

▶ Thread B

```
lock(mutex); /* access resource */  
mark resource free;  
signal(condition);  
unlock(mutex);
```

▶ Important to use while

Producer-consumer with mutex and condition variables

```
Producer() {  
    while(1){  
        produce data  
        lock(mutex)  
        while buffer_full()  
            c_wait(full,mutex)  
        insert item  
        if (counter == 1)  
            c_signal(empty)  
        unlock(mutex)  
    }  
}
```

```
Consumidor() {  
    while(1){  
        lock(mutex)  
        while buffer_empty()  
            c_wait(empty,mutex)  
        remove item  
        if (counter == size - 1)  
            c_signal(full);  
        unlock(mutex);  
        consume item  
    }  
}
```

- Important to wait on c.v. within critical section!
 - If lock after wait: race condition
- If wait did **not** unlock the mutex when going to sleep:
 - If wait after lock then deadlock

Producer-consumer with mutex and condition variables

```
Producer() {  
    while(1){  
        produce data  
        lock(mutex)  
  
        insert item  
        unlock(mutex)  
        c_signal(cond)  
    }  
}
```

```
Consumidor() {  
    while(1){  
        lock(mutex)  
        if buffer_empty()  
            c_wait(cond,mutex)  
  
        remove item  
  
        unlock(mutex);  
        consume item  
    }  
}
```

■ What is the problem here?

- May remove on an empty queue if producer interrupted before signal
- while instead of if!

UNIX operations for mutex

► Initialization

```
int pthread_mutex_init(pthread_mutex_t *mutex,  
                        pthread_mutexattr_t * attr);
```

► Destruction

```
int pthread_mutex_destroy(pthread_mutex_t *mutex);
```

► Lock a mutex; block the thread if lock taken

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

► Unlock a mutex

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

UNIX operations for condition variables

- ▶ Initialization

```
int pthread_cond_init(pthread_cond_t*cond,  
                      pthread_condattr_t*attr);
```

- ▶ Destruction

```
int pthread_cond_destroy(pthread_cond_t *cond);
```

- ▶ Suspend the thread and free associated mutex if wait on condition

- ▶ Until another thread signals that the condition holds
- ▶ When thread wakes up competes for mutex again

```
int pthread_cond_wait(pthread_cond_t*cond,  
                      pthread_mutex_t*mutex);
```

UNIX operations for condition variables

- ▶ Unblock threads and free mutex

```
int pthread_cond_signal(pthread_cond_t *cond);
```

- ▶ Signal threads blocked on condition
- ▶ No effect if no thread waiting (different from semaphores which increment regardless)

```
int pthread_cond_broadcast(pthread_cond_t *cond);
```

- ▶ All threads suspended on condition wake up

E.g.:

Producer-consumer with mutex (I)

```
#define MAX_BUFFER          1024          /* buffer size */
#define _DATA_      100000      /* data */

pthread_mutex_t mutex;          /* mutex to control access to buffer */
pthread_cond_t not_full;        /* condition variable */
pthread_cond_t not_empty;       /* condition variable */
int n_elements;                 /* nr elements in buffer */

int buffer[MAX_BUFFER];         /* shared buffer */

int main(int argc, char *argv[]){
    pthread_t th1, th2;

    pthread_mutex_init(&mutex, NULL);
    pthread_cond_init(&not_full, NULL);
    pthread_cond_init(&not_empty, NULL);
```

E.g.:

Producer-consumer with mutex (II)

```
pthread_create(&th1, NULL, Producer, NULL);
```

```
pthread_create(&th2, NULL, Consumer, NULL);
```

```
pthread_join(th1, NULL);
```

```
pthread_join(th2, NULL);
```

```
pthread_mutex_destroy(&mutex);
```

```
pthread_cond_destroy(&not_full);
```

```
pthread_cond_destroy(&not_empty);
```

```
return 0;
```

```
}
```


E.g.:

Producer-consumer with mutex (III)

```
void Producer(void) {
    int data, i ,pos = 0;

    for(i=0; i<_DATA_; i++ ) {
        data = i;                                /* produce data */
        pthread_mutex_lock(&mutex);              /* access the buffer */

        while (n_elements == MAX_BUFFER) /* if buffer full */
            pthread_cond_wait(&not_full, &mutex); /* block */

        buffer[pos] = i;

        pos = (pos + 1) % MAX_BUFFER;
        n_elements = n_elements + 1;

        if (n_elements == 1)
            pthread_cond_signal(&not_empty); /* buffer not empty */

        pthread_mutex_unlock(&mutex);
    }
    pthread_exit(0);
}
```

E.g.:

Producer-consumer with mutex (IV)

```
void Consumer(void) {
    int data, i ,pos = 0;

    for(i=0; i<_DATA_; i++ ) {
        pthread_mutex_lock(&mutex);          /* access the buffer */

        while (n_elements == 0)               /* if buffer empty */
            pthread_cond_wait(&not_empty, &mutex); /* block */

        data = buffer[pos];

        pos = (pos + 1) % MAX_BUFFER;
        n_elements = n_elements - 1 ;

        if (n_elements == MAX_BUFFER - 1);
            pthread_cond_signal(&not_full); /* buffer not full */

        pthread_mutex_unlock(&mutex);

        printf("Consume %d \n", data);      /* consume data */
    }
    pthread_exit(0);
}
```