

Message Passing



ARCOS Group

Distributed Systems

Bachelor In Informatics Engineering

Universidad Carlos III de Madrid

Basic concepts

- A mechanism for the communication and synchronization of processes executing on different machines
- “shared nothing” systems
- Defined as:
 - Set of processes having only local memory
 - Message passing different from a “call” since message contains content, not its address
 - Processes exchange messages – they must have a communication link
 - The transfer of data between processes requires cooperation (send has matching receive)
- Primitives:
 - `send(dest, message)`
 - `receive(src, message)`
- For connection-oriented communication also need to connect and disconnect



Distributed application paradigms - abstraction levels

Collaborative apps, object space


Services, object request broker, mobile agents

RPC, RMI


Client-server, P2P

Message passing


computer A




computer B




computer A




computer B



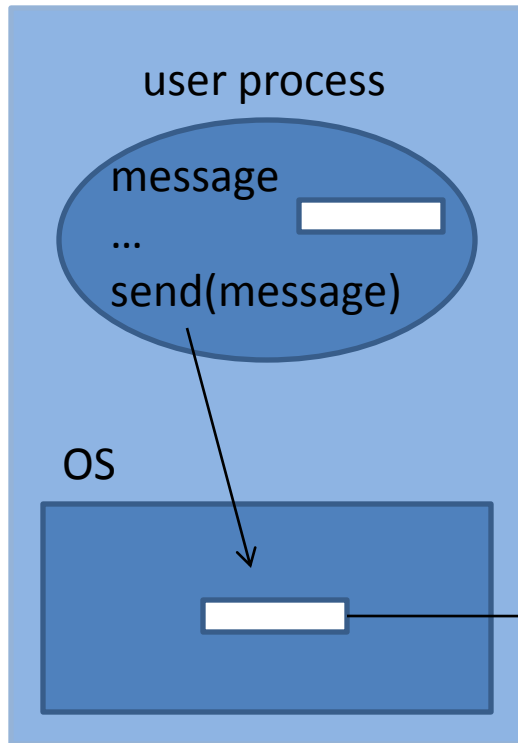
computer A





computer B



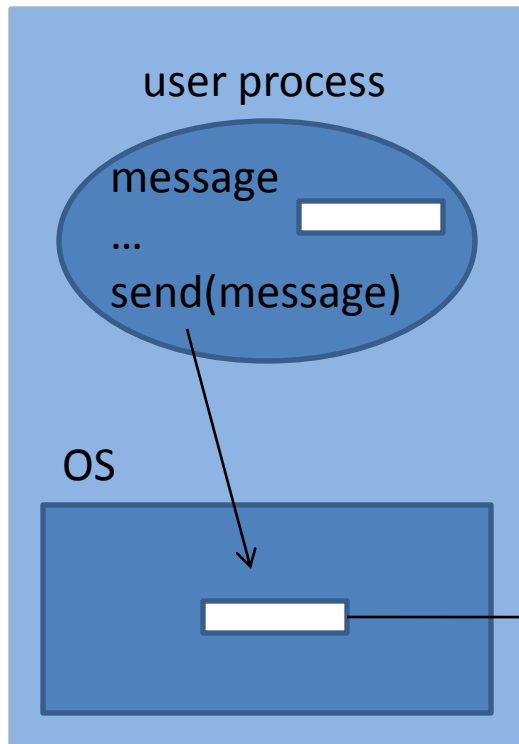
computer A





computer B



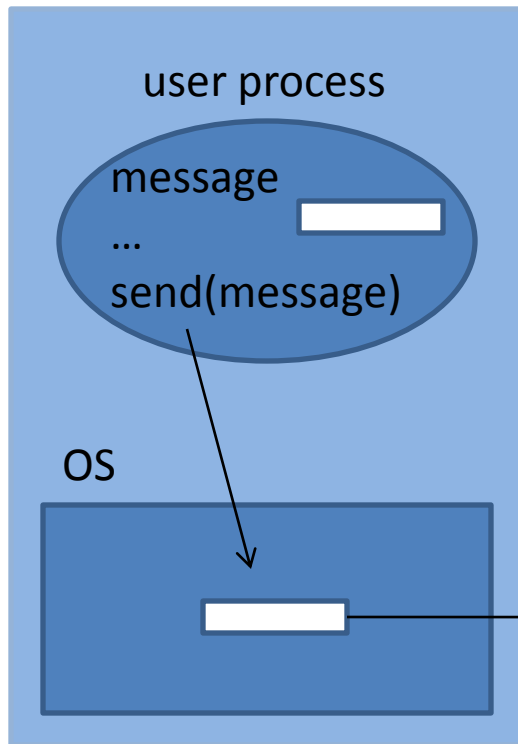
computer A



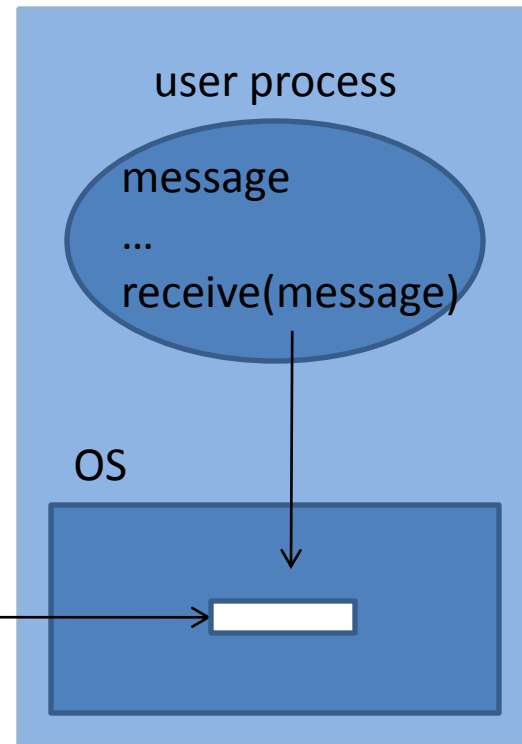
computer B



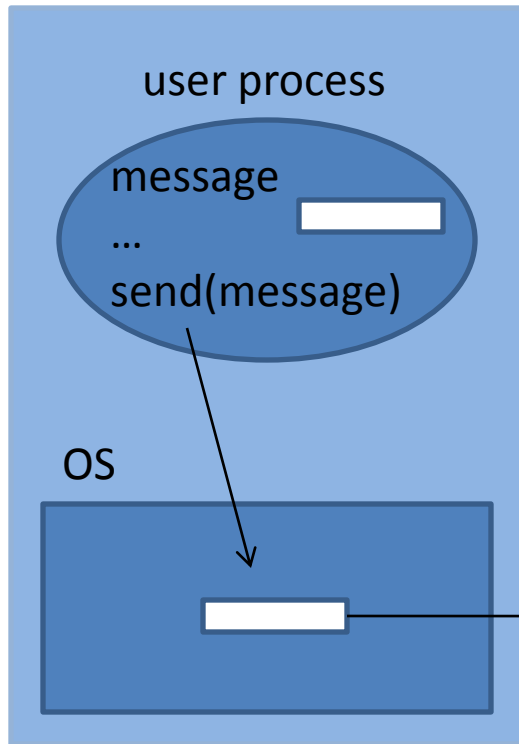
computer A



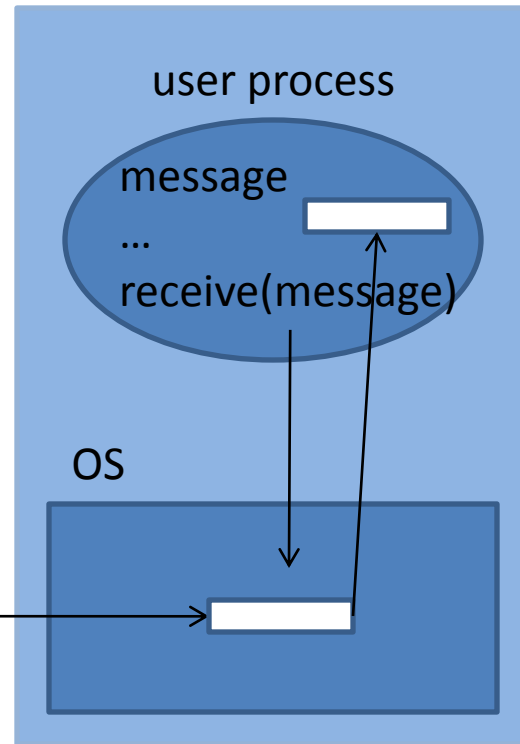
computer B



computer A



computer B



Other concurrency models besides message passing (MP)

- Data parallelism – determined by data partitioning
- Shared memory – common memory space
- Remote memory ops – process may access other process' memory w/o its participation
- Threads

Advantages and drawbacks of MP

- Advantages:
 - Hardware match – fits well on parallel supercomputers and clusters
 - Functionality – provides control not enabled by data parallelism and compiler-based models
 - Performance – explicit control of data locality
- Drawback:
 - Responsibility on the programmer: data distribution scheme, communication, synchronization w/o deadlock and race conditions

MP...



- ...may be implemented in the programming language, e.g. distributed object systems
- ...or in libraries, e.g. MPI
- Generally programming languages define messaging as asynchronous sending (by copy) of data item to communication endpoint
- MP forms the base of prominent theoretical foundations of concurrent computation

Design issues

- Message size: fixed vs variable
- Data flow: uni- vs bidirectional link
 - Can I both send and receive over a link?
- Naming message destination: direct vs indirect
 - Direct: specify the process to send to / receive from
 - Indirect: message queues, ports
- Communication type: synchronous vs asynchronous, blocking vs nonblocking
- Buffering: temporary copying of the message as part of the transmission protocol
- Reliability: are messages transferred reliably?
- Message order: are messages guaranteed to be delivered in order?

Naming

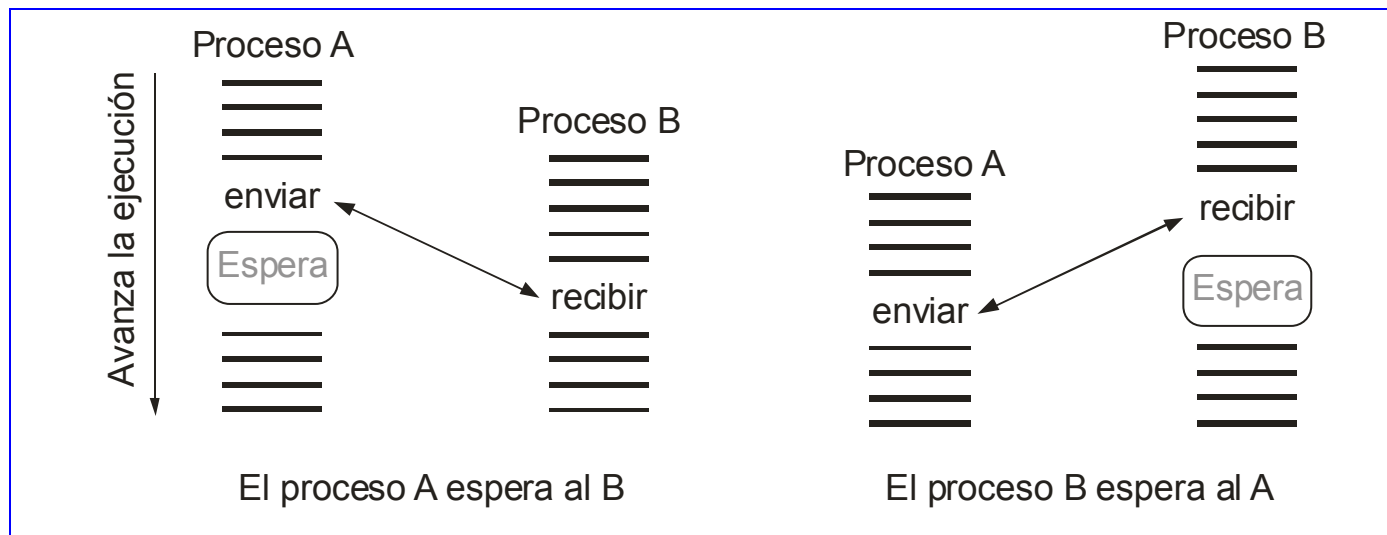
- Indirect: message sent to intermediary structures rather than to the process directly
 - Ports: associated with a process, e.g. sockets
 - Message queues: may be used by multiple senders / receivers, e.g. POSIX message queues



Communication type

- Synchronous: the sender does not return until matching receive in the destination process
 - Advantages: reasoning is simple, buffering not required (message stored on receiving side)
- Asynchronous: sender and receiver place no constraints on each other in terms of completion
 - Advantage: sender and receiver can overlap computation
 - Buffer may become full – must decide whether sender gets blocked (possible deadlock) or message is discarded (communication no longer reliable)
- Synchronous communication may be built on top of asynchronous by forcing the sender to wait for ACK from receiver before continuing

- Synchronous communication



Terminology

- Blocking communication: if the completion of the operation is dependent on certain events
 - For “send”: data successfully sent or safely copied s.t. buffer available for reuse
 - For “receive”: data must be safely stored in the receive buffer
- Nonblocking communication: operation returns w/o waiting for communication events to complete
- Synch/asynch implies blocking/nonblocking BUT
 - Not every blocking op is synchronous!
 - E.g. block on send till receiver machine has received the message, but receiver PROCESS may not have
 - Not every nonblocking is asynch!
 - E.g. synch nonblocking send completes if no receive_complete yet: send_start, send_complete (data copied out of send buffer)

- Solutions to unsafe programs

- Reorder operations

P0	P1
Send(P1)	Receive(P0)
Receive(P1)	Send(P0)

- Use operations which supplies receive buffer at the same time as send (if available)

P0	P1
Sendrecv(P1)	Sendrecv(P0)

- Use nonblocking operations

	P0	P1
return immediately	nbSend(P1)	Rnbeceive(P0)
	nbReceive(P1)	nbSend(P0)
ensure comm. completed	waitall	waitall

Buffering

- Refers to the feature of a communication protocol to temporarily copy the message into a buffer it controls
 - No buffering
 - With buffering: the intermediate structures (e.g. queues) have a given size
 - Queue not full: may send message
 - Otherwise: block or discard messages
 - If buffering is done in OS space it may lead to flooding!

UNIX messages queues

- Used for communication and synchronization
 - Use common queue to send/receive data
 - Blocking / nonblocking semantics
- Indirect naming
- Associated with files: processes can use queues only if they share the same file system
- Variable message size
- Bidirectional data flow
- Asynchronous send, synch/asynch receive
- Messages may have associated priorities

POSIX for managing message queues

- Create a queue with name and attribute

`mqd_t mq_open(char *name, int flag, mode_t mode, struct mq_attr *attr)`

flag:

`O_CREAT` Create queue if it doesn't exist

`O_RDONLY` Create a RO queue

`O_WRONLY` Create a WO queue

`O_RDWR` Create a RW queue

`O_NONBLOCK` nonblocking send and receive

mode: R, W, Exec rights

attr: max. nr of messages, message size, etc

- Close a queue

`int mq_close(mqd_t mqdes)`

- Delete a queue (after closing)

`int mq_unlink(char *name)`

E.g. create/destroy a message queue

```
#include <printf.h>          /* printf */
#include <fcntl.h>            /* flags */
#include <sys/stat.h>         /* modes */
#include <mqueue.h>
#define NUM_MSGS 50
int main () {
    mqd_t mqd;                /* Queue descriptor */
    struct mq_attr attributes; /* Attributes */

    attributes.mq_maxmsg = NUM_MSGS;
    attributes.mq_msgsize = sizeof(int);

    if ((mqd=mq_open("/store.txt", O_CREAT|O_WRONLY, 0777, &attributes))== -1){
        printf("Error mq_open\n");
        exit(-1);
    }
    mq_close(mqd);
    mq_unlink("/store.txt");
}
```

POSIX for managing message queues

- Modify queue attributes

```
int mq_setattr(mqd_t mqdes, struct mq_attr *qstat,  
               struct mq_attr *oldmqstat)
```

qstat: Contains new attributes

if oldmqstat != NULL it will store the new attributes

- Obtain queue attributes

```
int mq_getattr (mqd_t mqdes, struct mq_attr *qstat)
```


UNIX

- Send message to a queue

```
int mq_send(mqd_t mqdes, char *msg, size_t len,  
            int prio)
```

Message msg of size len sent to queue mqdes with priority prio

If queue full send may be blocking or not depending on **O_NONBLOCK**
and it returns nr of bytes sent or -1 (error)

- Receive message from a queue

```
int mq_receive(mqd_t mqdes, char *msg,  
               size_t len, int *prio)
```

Extract message with highest priority

If queue empty block receive if not **O_NONBLOCK**
and return nr of bytes read or -1 (error)

Producer-consumer with MP

```
Producer () {  
    for(;;) {  
        <Produce data>  
        send(Consumer, data);  
    } /* end for */  
}
```

```
Consumer () {  
    for(;;) {  
        receive(Producer, data);  
        <Consume data>  
    } /* end for */  
}
```

Producer-consumer with message queues

```
#include <mqueue.h>
#include <stdio.h>
#define MAX_BUFFER      1024      /* buffer size */
#define PROD_DATA      100000    /* data produced */

mqd_t store;                  /* queue for inserting produced
                              data and removing consumed data */

void main(void) {
    struct mq_attr attr;
    attr.mq_maxmsg = MAX_BUFFER;
    attr.mq_msgsize = sizeof (int);
    store = mq_open("STORE", O_CREAT|O_WRONLY, 0700, &attr);
    if (store == -1) {
        perror ("mq_open");
        exit(-1);
    }
    Producer();
    mq_close(store);
    exit(0);
}
```

```
Producer(void){
    int data;
    int i;

    for(i=0;i<PROD_DATA;i++){
        /* produce data */
        data = i;
        if (mq_send(store, &data, sizeof(int), 0)== -1){
            perror("mq_send");
            mq_close(store);
            exit(1);
        }
    } /* end for */
    return;
} /* end producer */
```

```
#include <mqueue.h>
#include <stdio.h>
#define MAX_BUFFER      1024      /* buffer size */
#define PROD_DATA      100000    /* data produced */

mqd_t store;

void main(void) {
    struct mq_attr attr;
    store = mq_open("STORE", O_RDONLY);
    if (store == -1) {
        perror ("mq_open");
        exit(-1);
    }
    Consumer();
    mq_close(store);

    exit(0);
}
```

```
Consumer(void){
    int data;
    int i;

    for(i=0;i<PROD_DATA;i++){
        /* receive data */
        data = i;
        if (mq_receive(store,
                        &data, sizeof(int), 0)== -1){
            perror("mq_send");
            mq_close(store);
            exit(1);
        }
        /* Consume data */
        printf("data consumed: %d\n", data);
    } /* end for */
    return;
} /* end consumer */
```

Group communication

- Data movement
 - One-to-many: multicast, broadcast
 - Many-to-one: client-server
 - All-to-all
 - Gather/scatter
- Synchronization
- Collective computation (reductions)
- Useful for:
 - Fault tolerance via replication of services
 - Better performance via data replication
 - Multiple updates
 - Reduce operations for parallel computation
- Collective operations are blocking

Examples of MP-style languages

- Actors
 - Amorphous computing
 - Antiobjects
 - Flow-based programming
 - SOAP protocol
 - Smalltalk
-
- Lambda calculus is considered the earliest MP programming language

Actors

- Before actors concurrency defined in terms of threads, locks, and buffers
- Characteristics:
 - Asynchronous communication
 - Direct naming: message recipients identified by address
 - Recipient address may be included in message
 - No restriction on message arrival – similar to packet switching systems, enables optimizations: buffering packets, send on different paths ,resend, pipeline message processing, etc
 - In most formalisms message delivery is guaranteed
 - Not necessarily buffered

MPL programming: performance guidelines

- Start with tuned serial program
- Control process granularity: increased granularity decreases communication /computation overhead
- Overlap communication and computation: use nonblocking communication to hide comm. costs
- Avoid unnecessary synchronization
- Avoid buffering when possible: extra copies of messages decrease bandwidth
- Keep data local