

Client/Server paradigm



ARCOS Group

Distributed Systems

Bachelor In Informatics Engineering

Universidad Carlos III de Madrid

A simple definition


“ Server software accepts requests for data from client software and returns the results to the client”

Where are operations executed

- Most of the application processing is done on a computer (client side)
- ...which obtains application services (such as database services) from another computer (server side) in a master slave configuration

Client/Server paradigm

- Assigns different roles to the communicating processes
- Server:
 - Offers services
 - Passive: waits for incoming requests from clients
- Client:
 - Requests services
 - Active: sends requests to server(s)



Servers may be...

- Depending on the type of connection with the client:
 - Connection-oriented
 - Connectionless
- Depending on the number of serviced client sessions:
 - Sequential: if it communicates with a single client session at the time
 - Concurrent: if it may communicate with multiple client sessions at the time
- Depending on whether it stores communication state:
 - Stateful
 - Stateless

Connection-oriented servers

- Client and server must establish a connection (logical or physical) before communicating; when finalizing the communication they must close the connection
- Once the connection established it isn't necessary to refer directly to the sender and receiver
- C-O protocols may be seen as stateful – keep track of conversation
- May not serve new client before current one closes the connection!
- E.g.: TCP

Connectionless servers


- Data exchanged via self-contained packets which must contain explicit server/client address information
 - No previous agreement
- C-less protocols may be seen as stateless
- May interleave different client requests!
- E.g.: IP, UDP

What is a session?

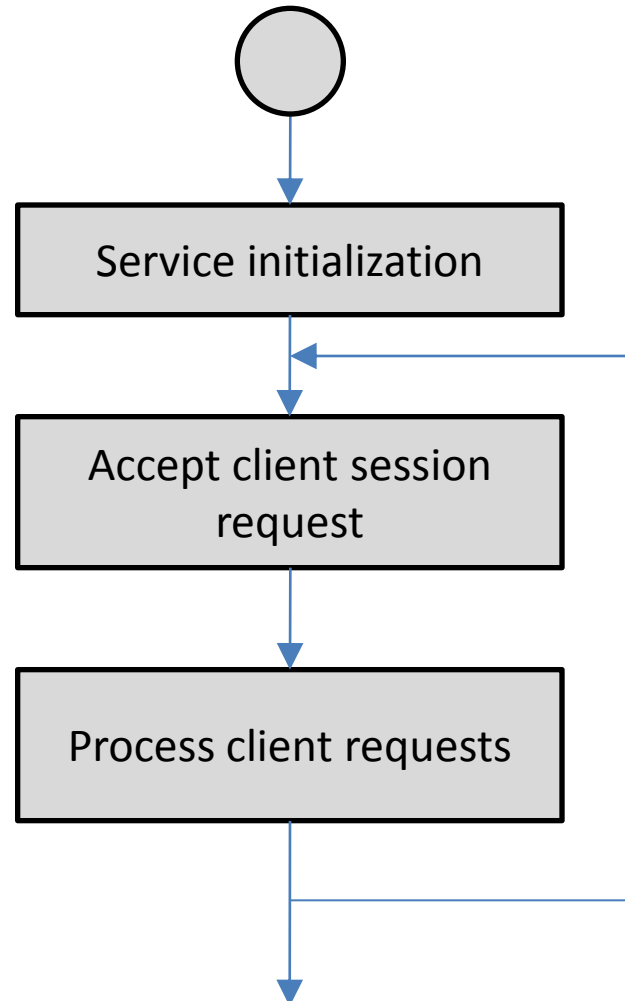
- ▶ Session: Interaction between client and server until client gets the requested service
- ▶ The server runs an infinit loops which accepts service requests from client sessions
- A **service protocol** specifies the rules that the client and server follow during a session wrt:
 - Naming a service: services identify themselves via a registered logical name or the server physical process address (machine name + port number)
 - Communication sequence
 - Data representation

Sequential servers

- Serves client requests sequentially
- Does not interleave requests from multiple client sessions
- While attending to a client session all others must wait

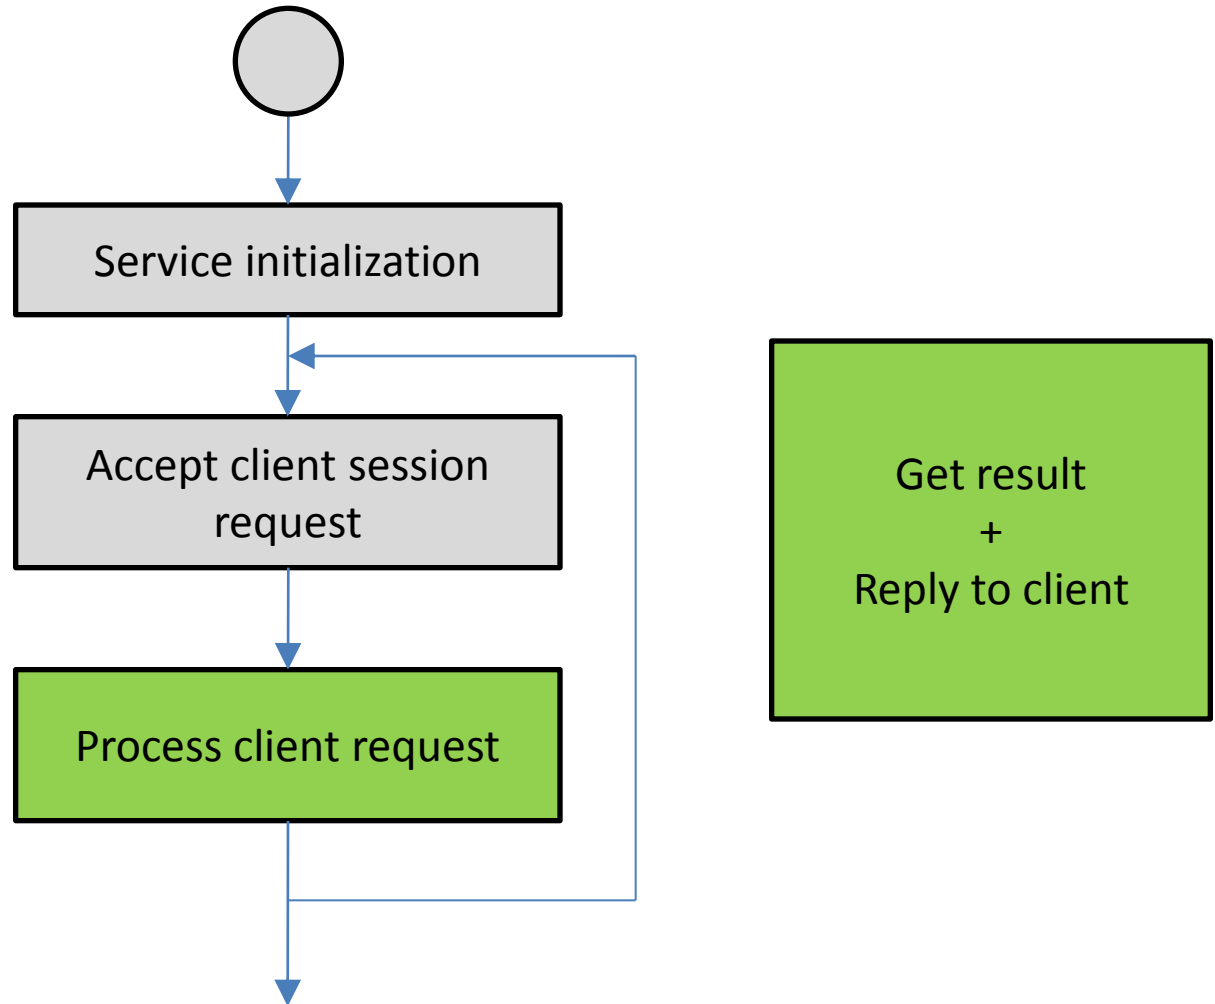


Execution flow



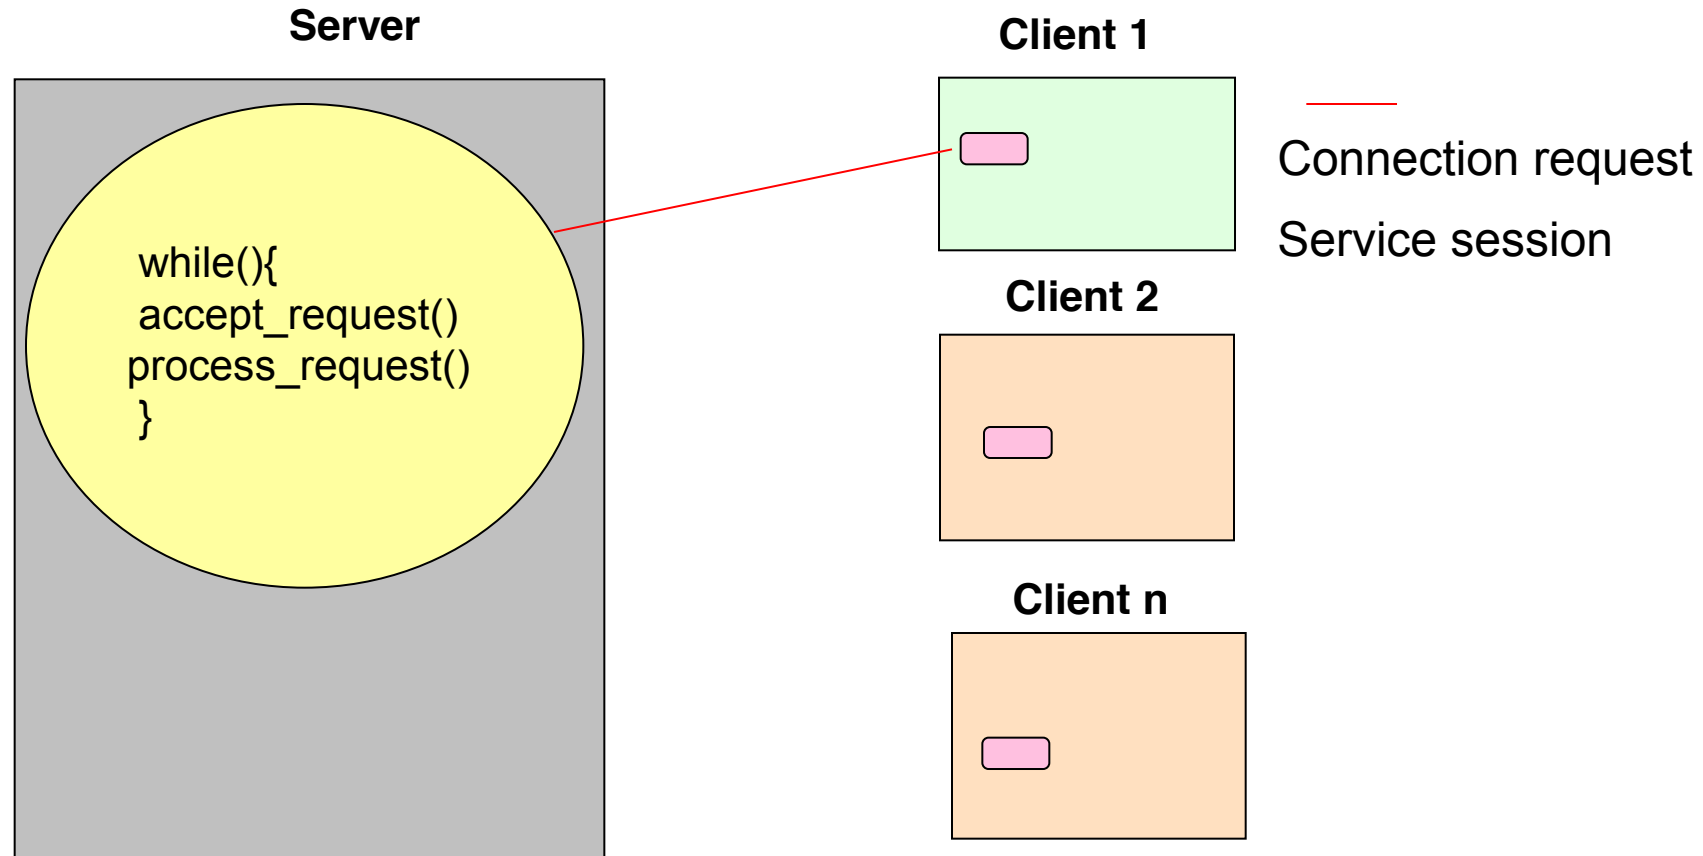
If for each client i keep a service session this server is sequential!

Execution flow

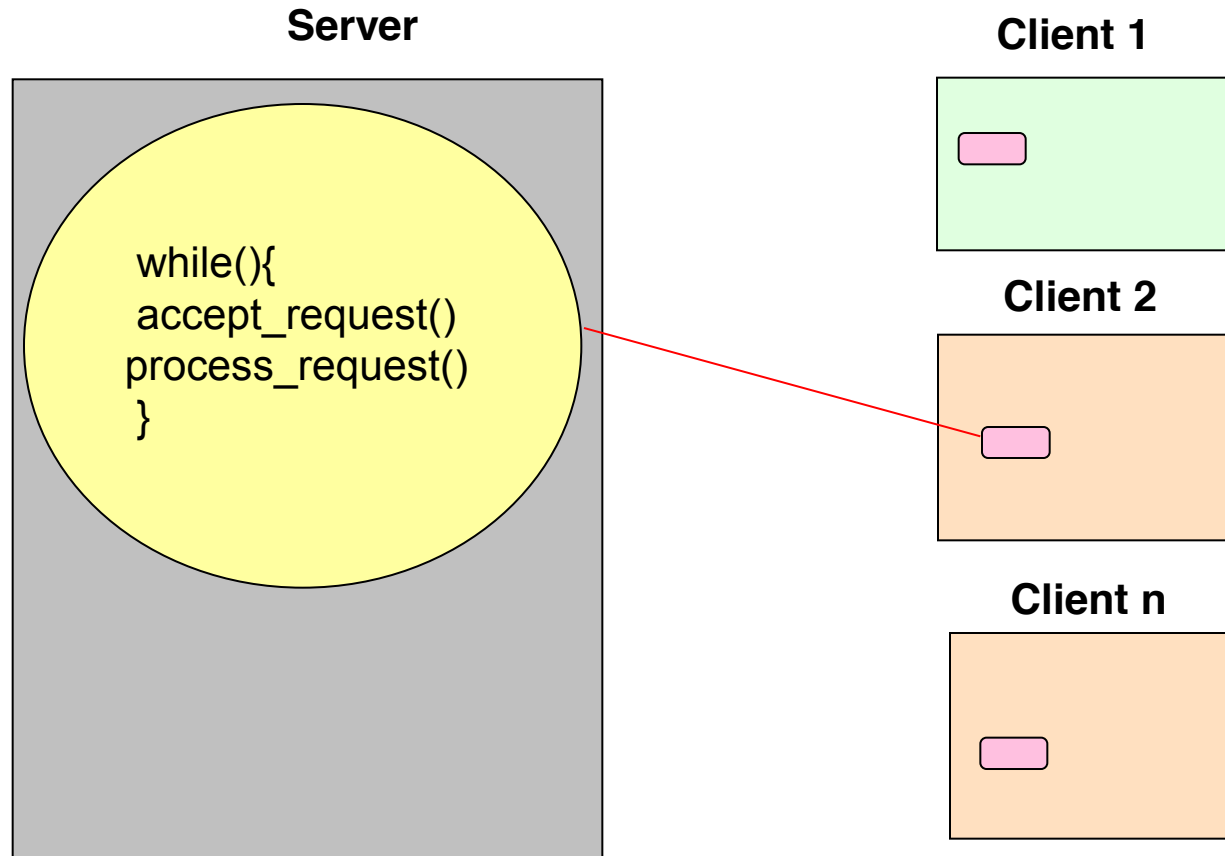


Sequential server

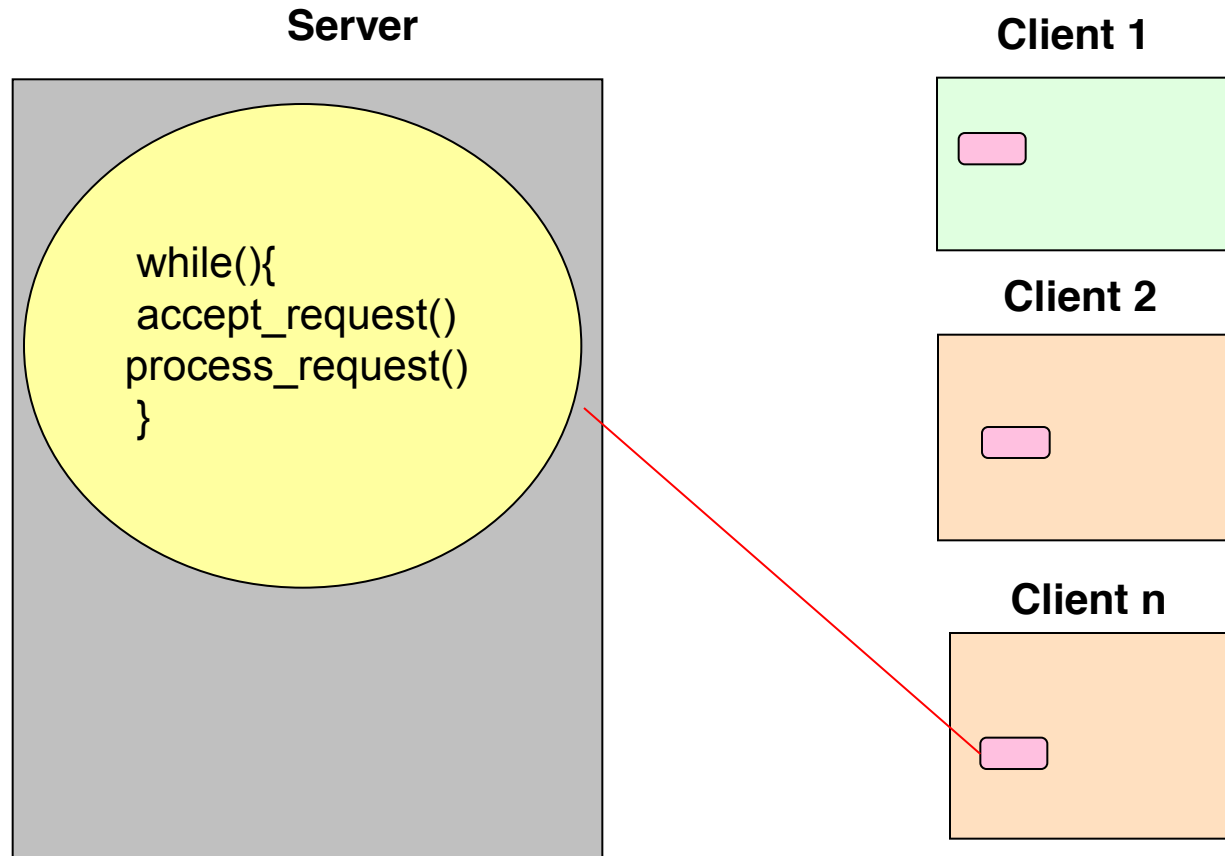
Sequential Client/Server



Sequential Client/Server

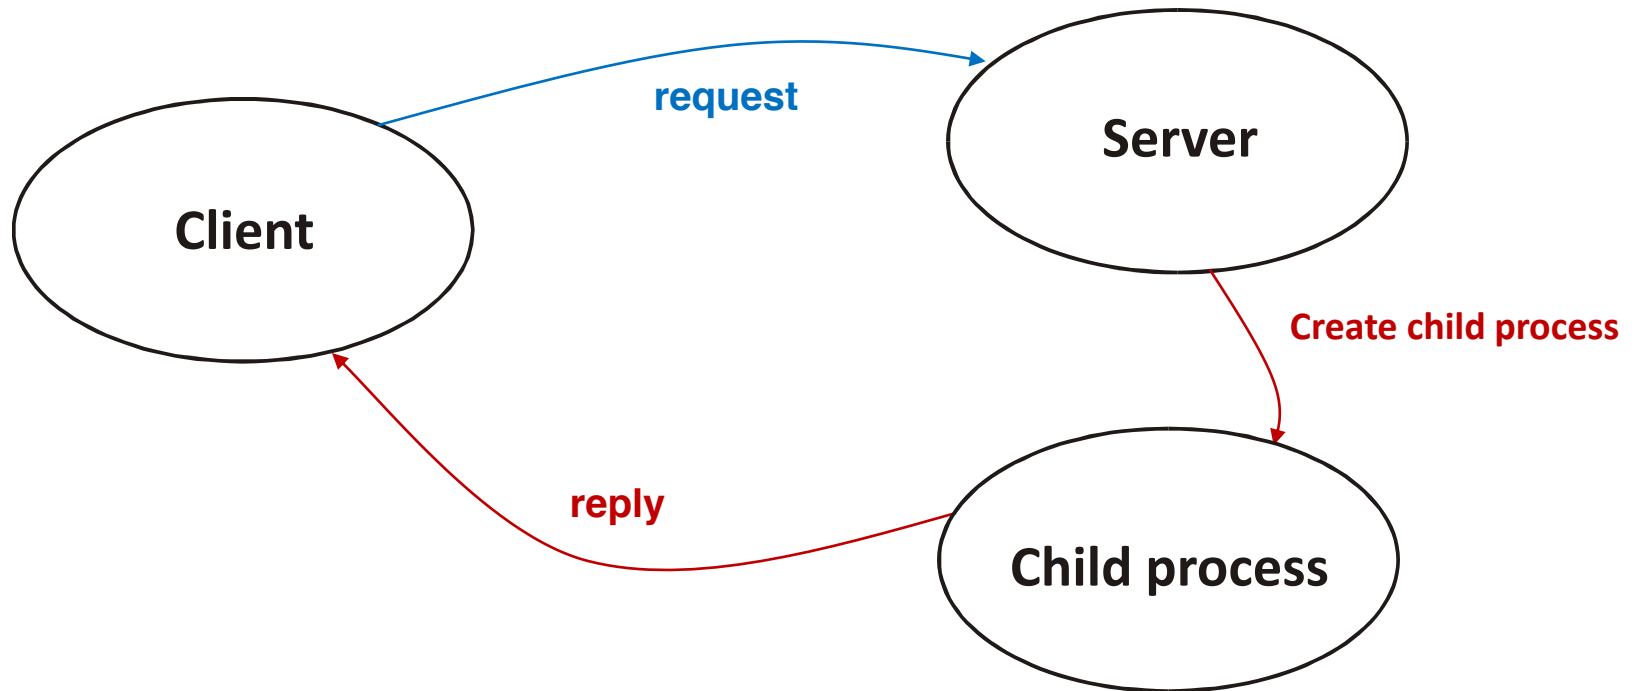


Sequential Client/Server



Concurrent servers

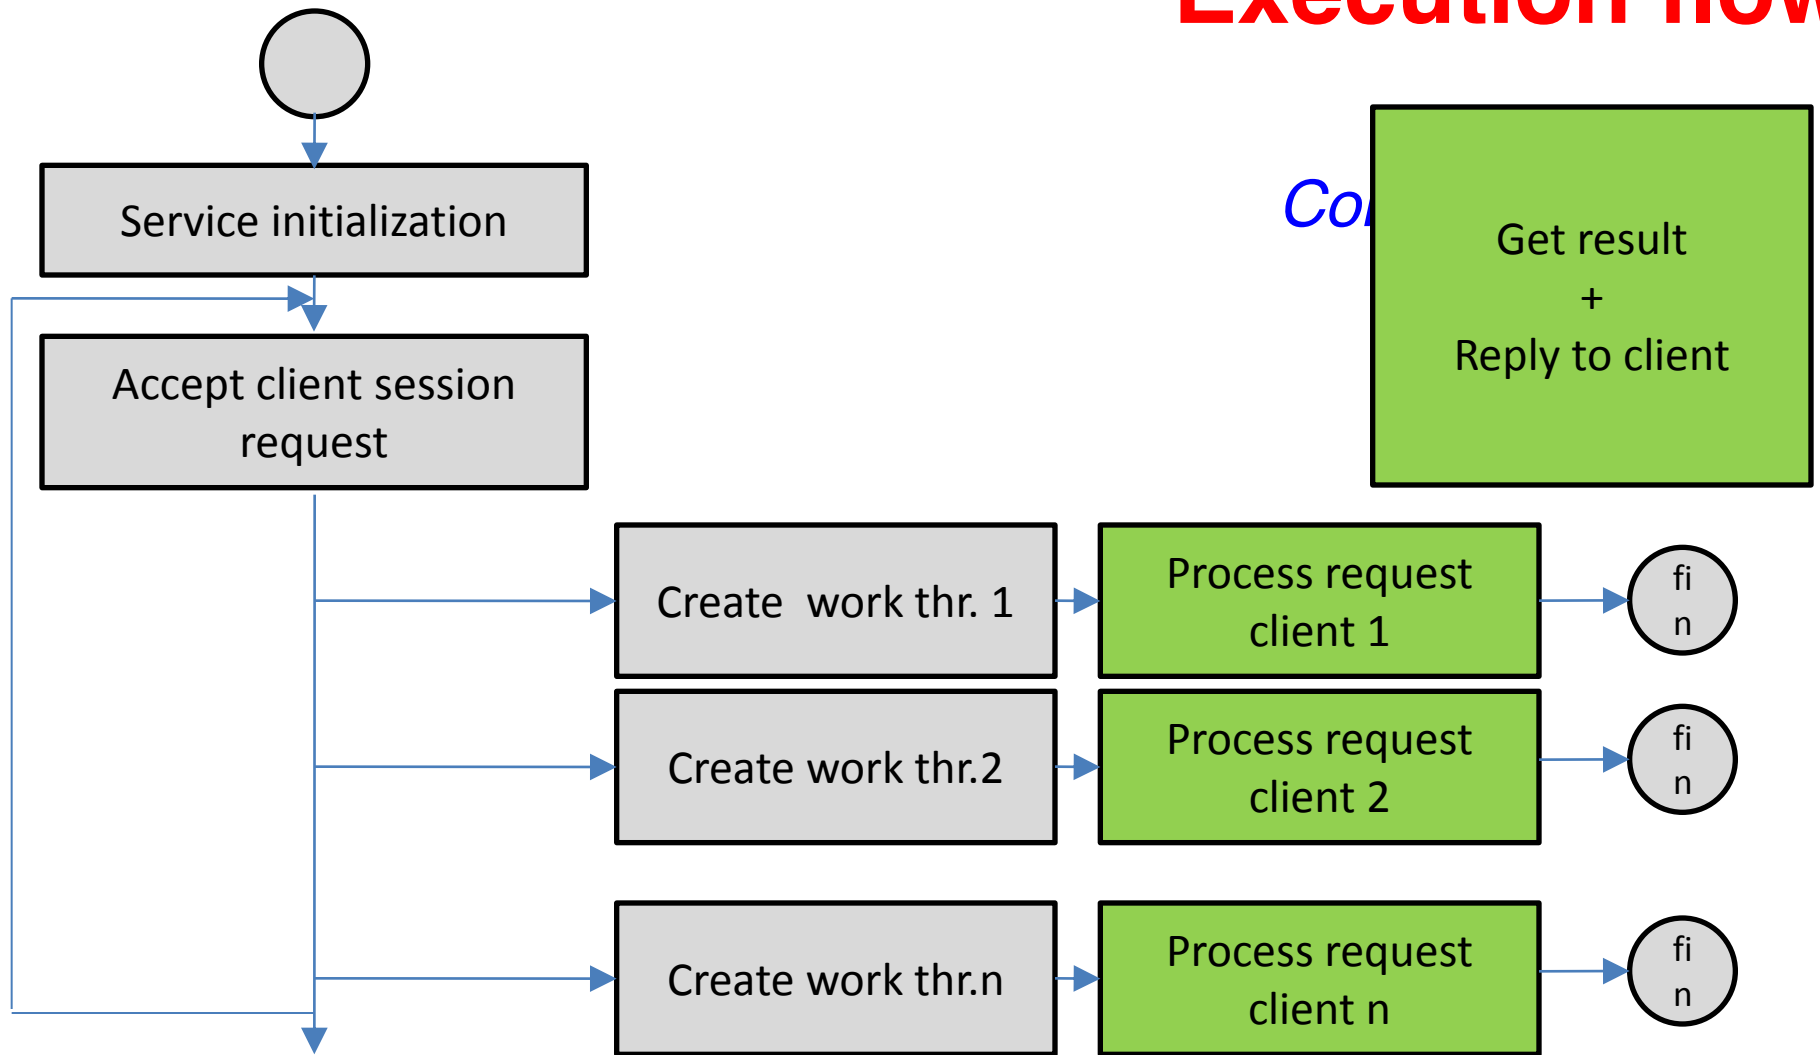
- Server creates a child process which will process the request and send the reply to the client
- Multiple client sessions may be interleaved



Concurrent server

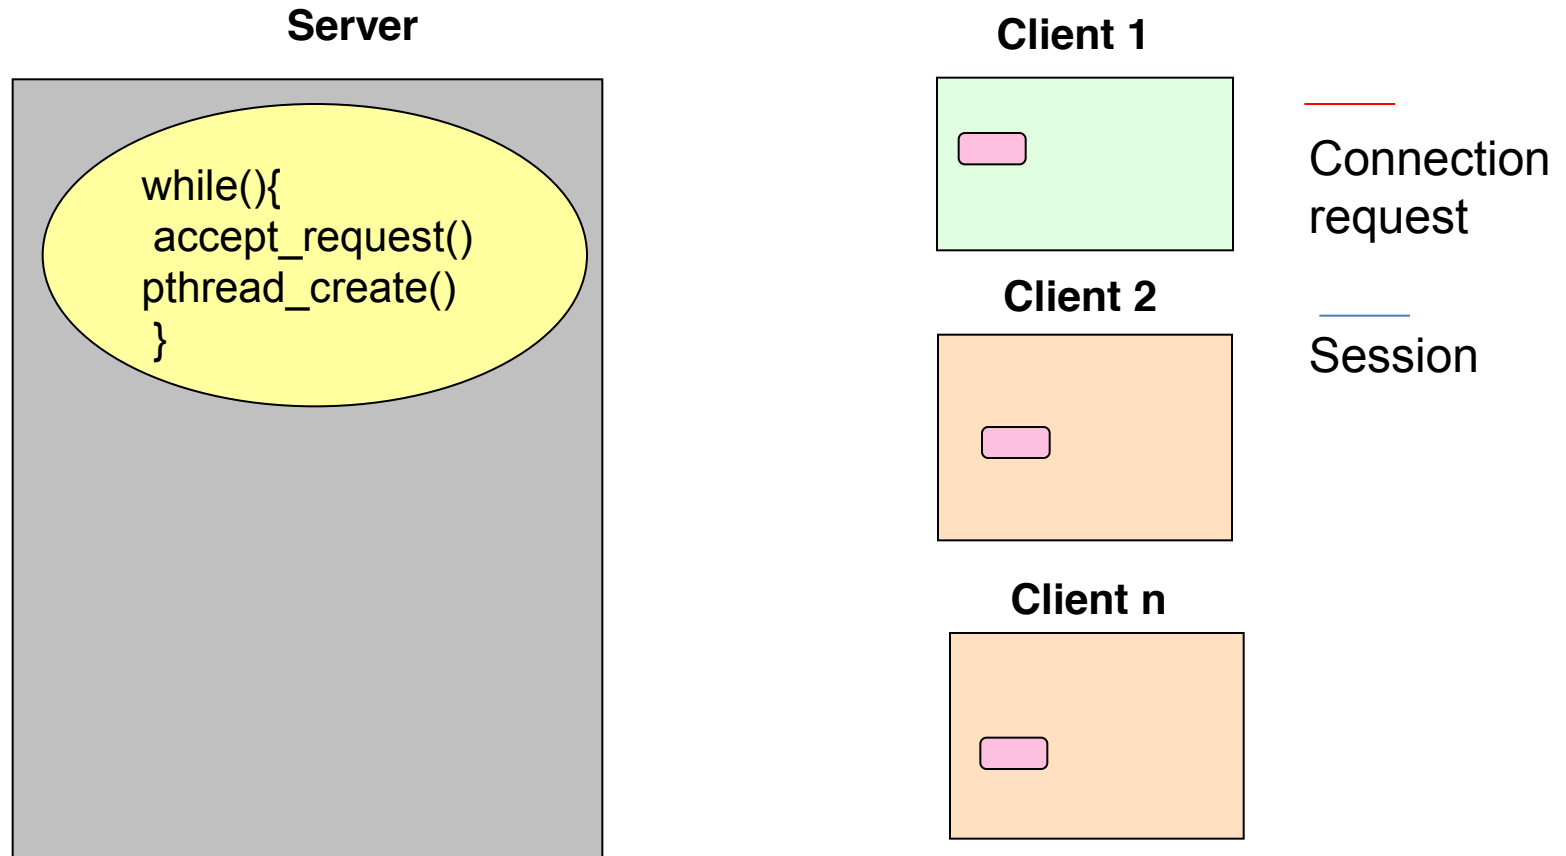


Execution flow

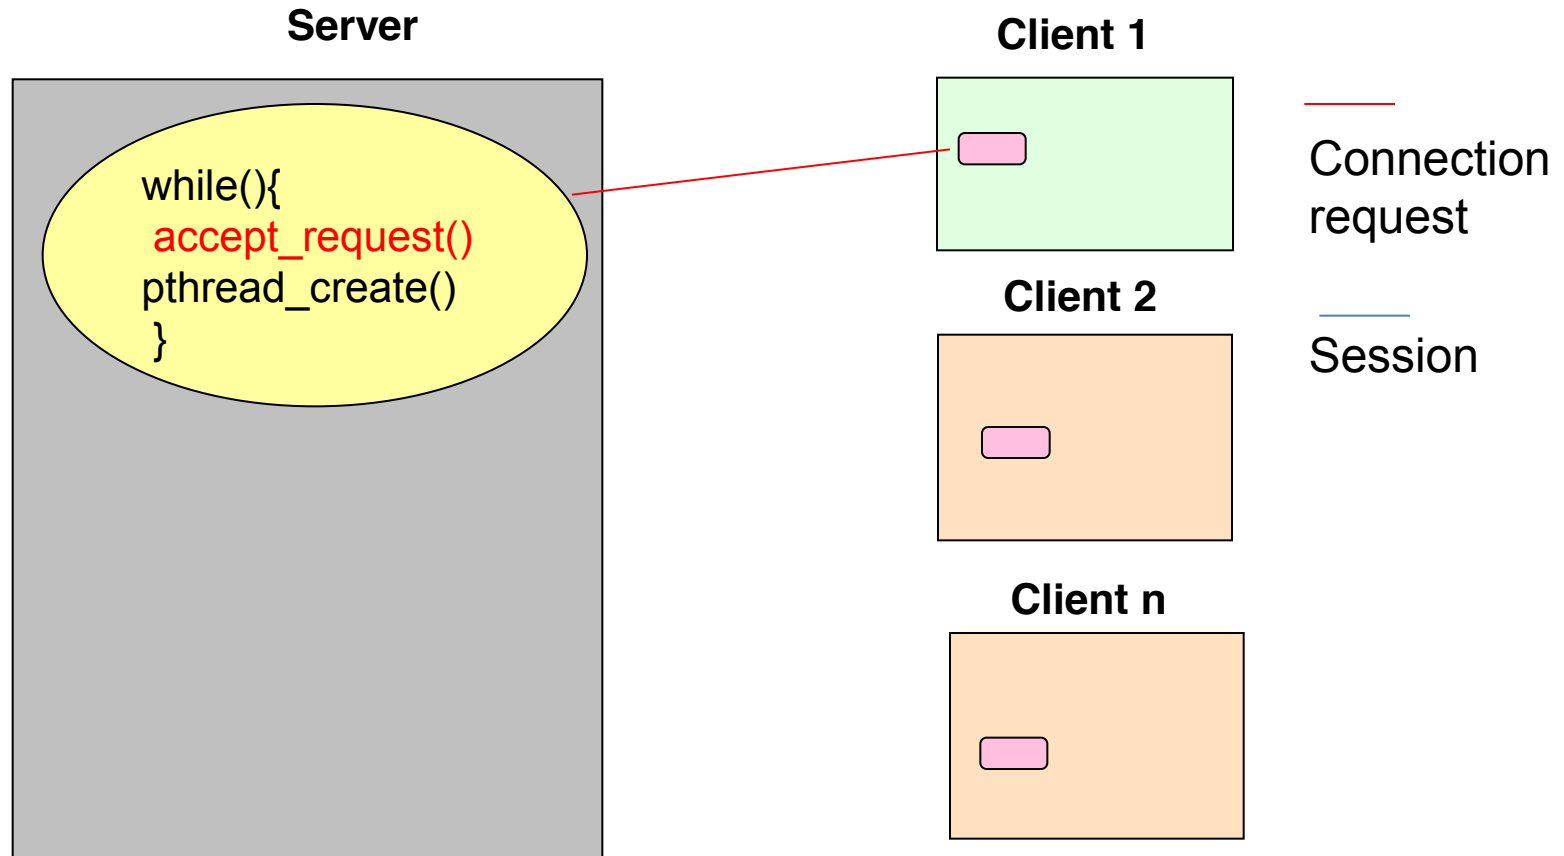


May use asynchronous inter process communication primitives instead of threads!

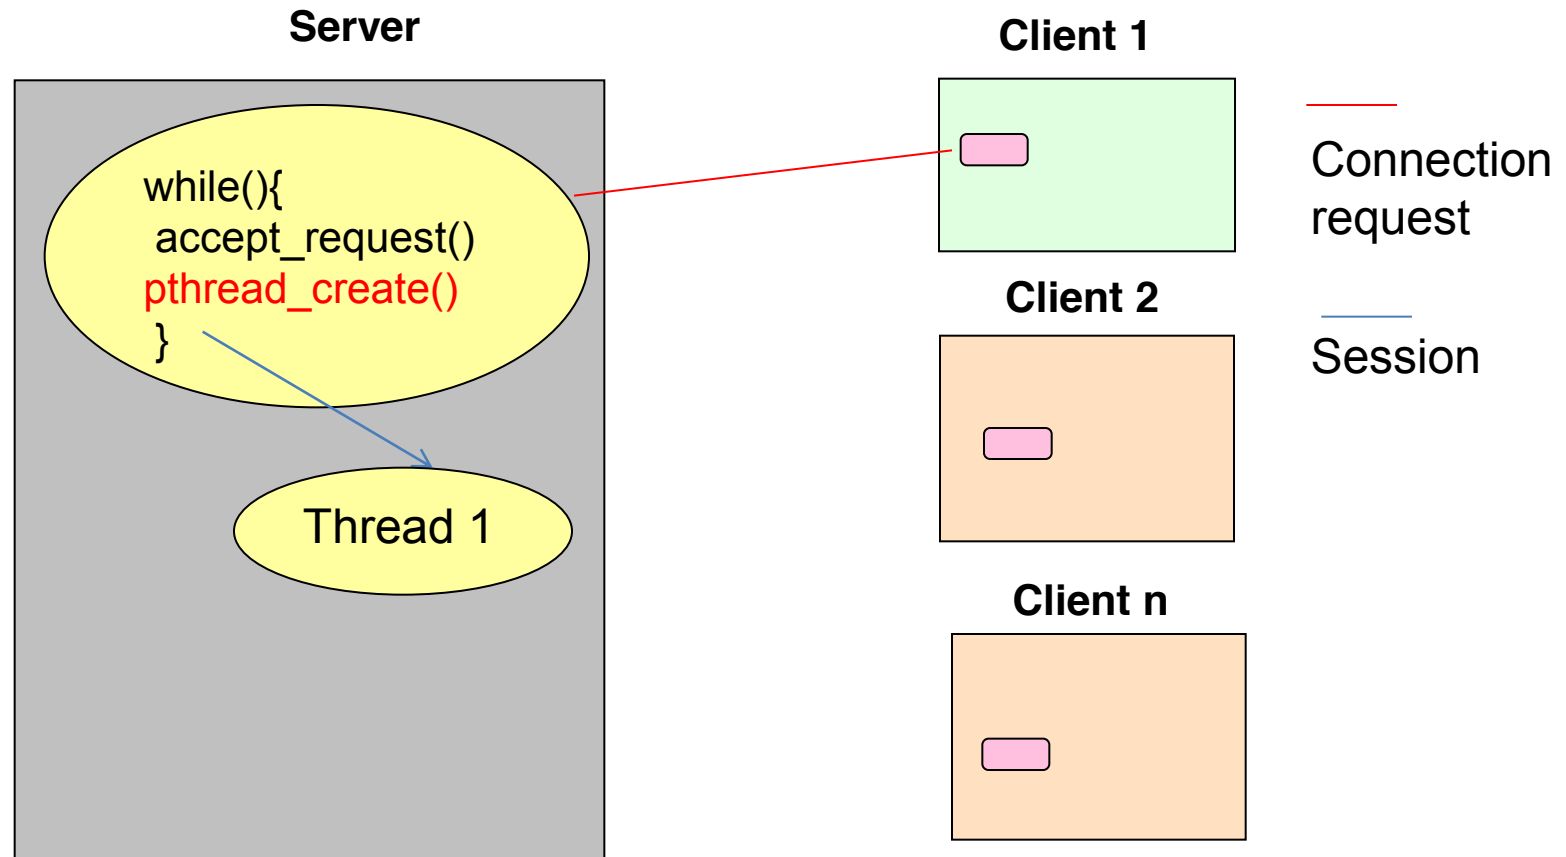
Concurrent Client/Server



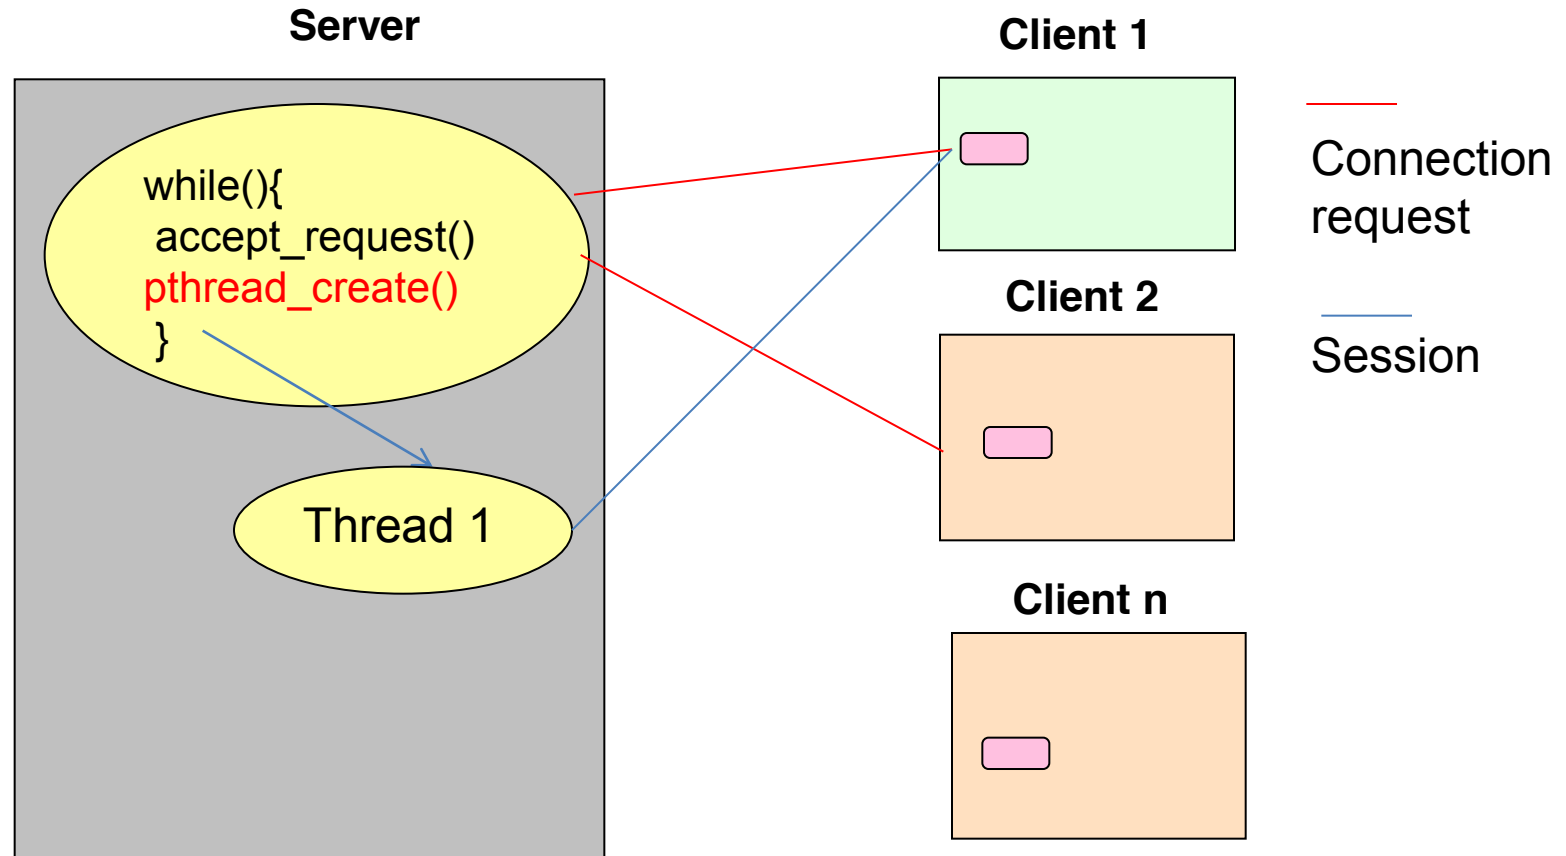
Concurrent Client/Server



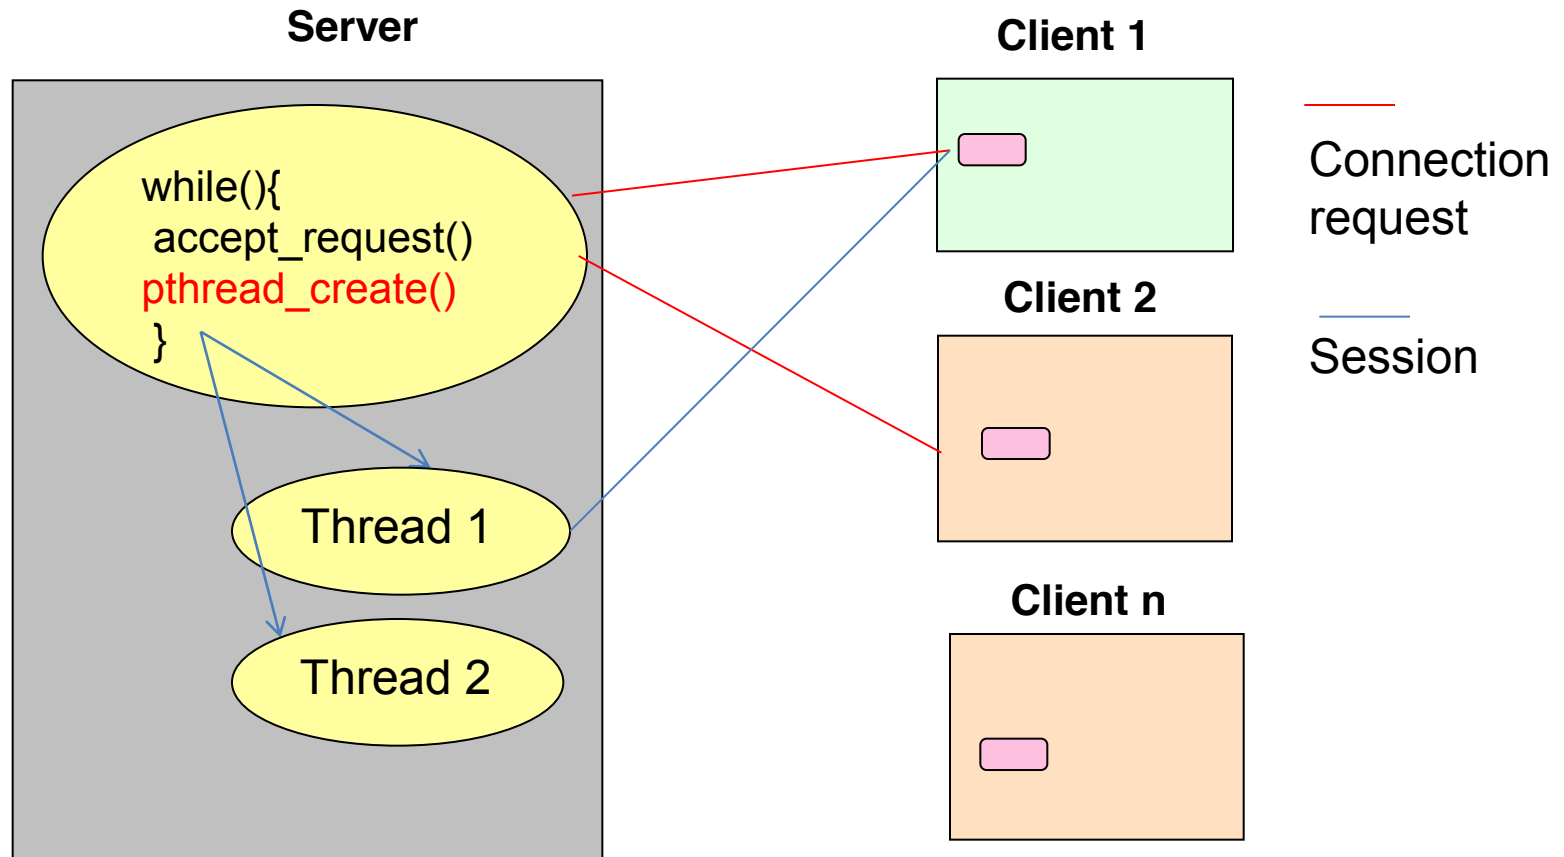
Concurrent Client/Server



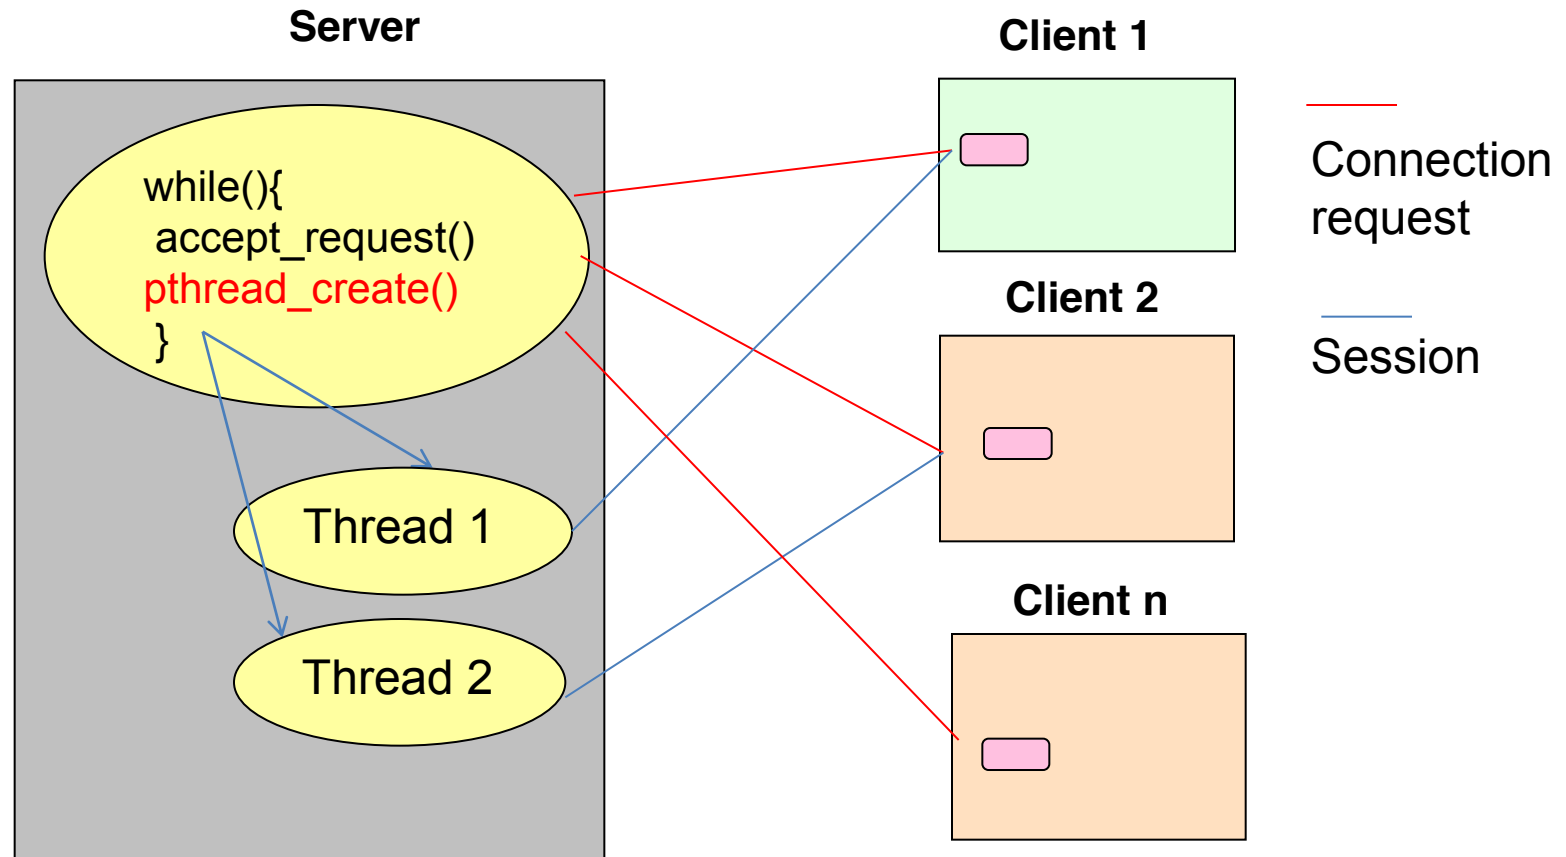
Concurrent Client/Server



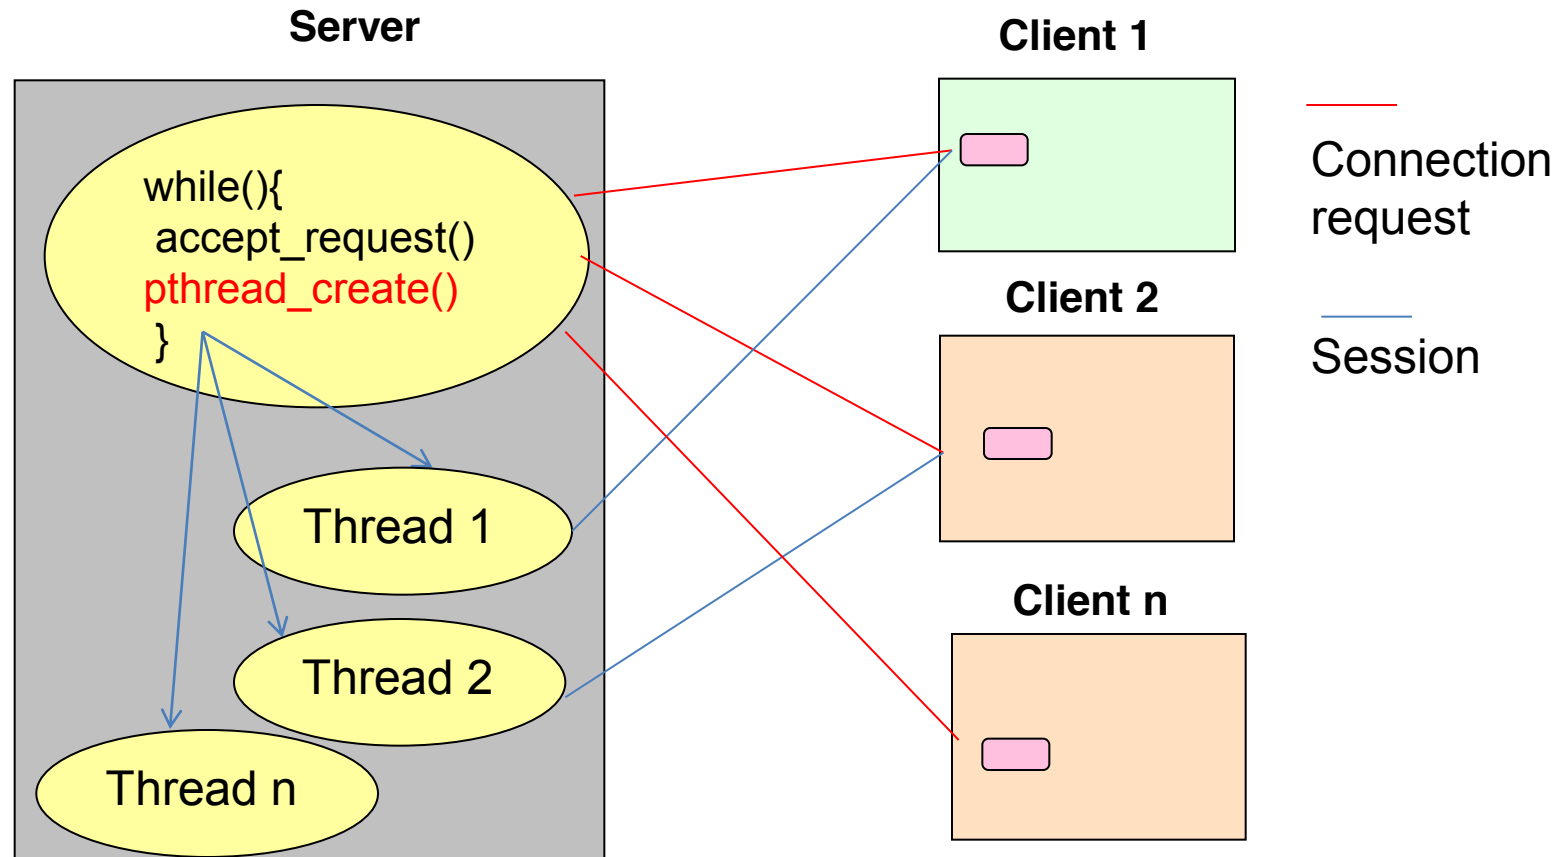
Concurrent Client/Server



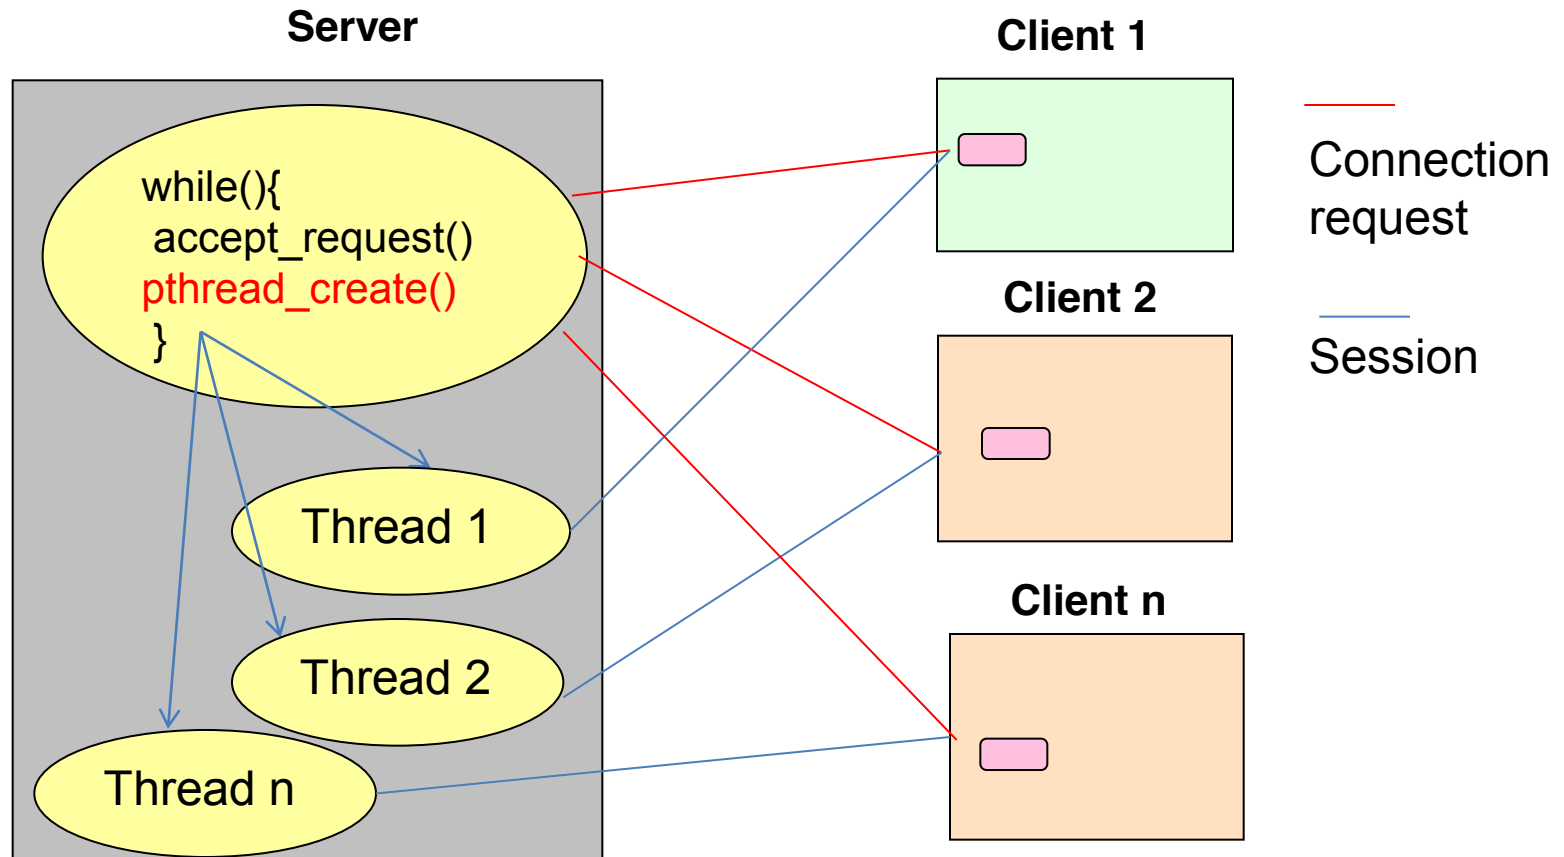
Concurrent Client/Server



Concurrent Client/Server




Concurrent Client/Server



Server design using *threads*

- ▶ Several alternatives to build parallel servers:
 - ▶ A single process which accepts requests and either (1) distributes them to threads from a thread pool or (2) creates a new thread to service the request
 - ▶ Set of similar threads which can read requests from a port
 - ▶ Pipeline the work and have a specialized thread for each stage



Servers may be...

- Stateless:
 - Every request and reply is independent
 - No state maintained by the server
 - Client may maintain session state and send it as part of the service request to the server
 - Client: "Send me block 1 of file "xxx" from directory "dir"
 - Server: "Here it is"
 - <more of the same>
 - E.g.: [HTTP](#)
- Stateful:
 - Maintains state information
 - Each request/reply may depend on previous ones
 - Client: "Send me file "xxx" from directory "dir"
 - Server: "Here is block 0 of file "xxx"
 - Client: "I have it"
 - Server: "Here is block 1 of file "xxx"
 - E.g.: [Telnet](#)

Stateful servers

- Global state:
 - Information common to all clients
 - E.g.: “time of day” server
- Session information
 - Information specific to each client session
 - E.g.: FTP (File Transfer Protocol)


Clients may be...

- Thin: also called lean client or slim client
 - Depends heavily on other computer (e.g. its server) to process most or all of its business logic - which traditional systems like fat clients take on
 - E.g. : the server may need to provide data persistence, process information on client's behalf, etc
 - May be seen as amortizing computing services across several user interfaces
 - Problem: server become single point of failure!
 - Good for checking security thread models
 - Bad if denial of service attack from a client
- Thick: also called fat client or heavy client


SW architecture

- Usually consists of three layers/tiers:
 - Presentation: user interface issues
 - Application logic: isolates data processing in one location and maximizes reuse, modification in services does not affect presentation
 - Server needs to process client request, compute result and return it to client
 - Client needs to send service request and visualise result
 - Services - we need two types:
 - On server - those processing the request
 - Some IPC mechanism!
 - Must be able to manage data
- May seem similar but different from MVC architecture!
 - View sends updates to controller; controller updates the model, view gets updated directly from model
 - Model = data +domain logic (+persistence, notification)
 - View = query model, render view
 - Controller = init model, wiring up events between controller and V/M


Client (dumb) - Server Model



True Client-Server Model



Distributed Client-Server Model




Typical to fat clients

If too complex may want to split into a three-tier

Client/Server applications using message queues


- Distributer thread:
 - Each request results in creating a new work thread which:
 - Processes the request
 - Sends reply to client
 - When client session finishes the thread is destroyed
- Concurrent model:
 - Distributer and work threads execute concurrently

Example: add two numbers




Example: add two numbers

Machine A





$\text{add}(5,2)$
Result = 7

Machine B



Machine C




Example: data types

```
#define MAXSIZE          256

struct request {
    int a;                /* op. 1 */
    int b;                /* op. 2 */
    char q_name[MAXSIZE]; /* client queue name - this is
                           where the server sends the reply
                           to */
};
```

The structure of a multithread server



Multithread server with message queues (I)

```
#include "mensaje.h"
#include <mqueue.h>
#include <pthread.h>
#include <stdio.h>

/* mutex and condition variables for the message copy */
pthread_mutex_t mutex_msg;
int msg_not_copied = TRUE;          /* TRUE = 1 */
pthread_cond_t cond_msg;

int main(void)
{
    mqd_t q_server;                  /* server queue */
    struct request msg;               /* message to receive */
    struct mq_attr q_attr;           /* queue attributes */
    pthread_attr_t t_attr;           /* thread attributes */

    q_attr.mq_maxmsg = 20;
    q_attr.mq_msgsize = sizeof(struct request);
```

Multithread server with message queues (II)

```
q_server = mq_open("SERVER", O_CREAT|O_RDONLY, 0700, &attr);
if (q_server == -1) {
    perror("Can't create server queue");
    return 1;
}

pthread_mutex_init(&mutex_msg, NULL);
pthread_cond_init(&cond_msg, NULL);
pthread_attr_init(&attr);

/* thread attributes */
pthread_attr_setdetachstate(&t_attr, PTHREAD_CREATE_DETACHED);
```


Multithread server with message queues (III)

```
while (TRUE) {  
    mq_receive(q_server, &msg, sizeof(struct request), 0);  
  
    pthread_create(&thid, &attr, process_message, &msg);  
  
}  
}
```

Multithread server with message queues (IV)

```
while (TRUE) {  
    mq_receive(q_server, &msg, sizeof(struct request), 0);  
  
    pthread_create(&thid, &attr, process_message, &msg);
```

⋮

Race condition!!

```
}  
}
```

Multithread server with message queues (V)

```
while (TRUE) {  
    mq_receive(q_server, &msg, sizeof(struct request), 0);  
  
    pthread_create(&thid, &attr, process_message, &msg);  
  
    /* wait for thread to copy message */  
    pthread_mutex_lock(&mutex_msg);  
    while (message_not_copied)  
        pthread_cond_wait(&cond_msg, &mutex_msg);  
    message_not_copied = TRUE;  
    pthread_mutex_unlock(&mutex_msg);  
  
} /* FIN while */  
} /* Fin main */
```

Critical section

Multithread server with message queues (VI)

```
void process_message(struct mensaje *msg){
    struct request msg_local;           /* local message */
    struct mqd_t q_client;              /* client queue */
    int result;

    /* thread copies message to local message*/
    pthread_mutex_lock(&mutex_msg);
    memcpy((char *) &msg_local, (char *)&msg, sizeof(struct
        request));

    /* wake up server */
    message_not_copied = FALSE;         /* FALSE = 0 */

    pthread_cond_signal(&cond_msg);

    pthread_mutex_unlock(&mutex_msg);
```

Multithread server with message queues (VII)

```
/* execute client request and prepare reply */
result = msg_local.a + msg_local.b;

/* return result to client by sending it to queue */
q_client = mq_open(msg_local.name, O_WRONLY);

if (q_client == -1)
    perror("Can't open client queue */
else {
    mq_send(q_client, (char *) &result, sizeof(int), 0);
    mq_close(q_client);
}
pthread_exit(0);
}
```

Client process

```
#include "mensaje.h"
#include <mqueue.h>
void main(void) {
    mqd_t q_server;          /* server message queue */
    mqd_t q_client;         /* client message queue */

    struct request req;
    int res;
    struct mq_attr attr;

    attr.mq_maxmsg = 1;
    attr.mq_msgsize = sizeof(int);
    q_client = mq_open("CLIENT_ONE", O_CREAT|O_RDONLY, 0700, &attr);

    q_server = mq_open("ADD_SERVER", O_WRONLY);

    /* fill in request */
    req.a = 5;      req.b = 2;      strcpy(req.q_name, "CLIENT_ONE");

    mq_send(q_server, &req, sizeof(struct request), 0);
    mq_receive(q_client, &res, sizeof(int), 0);

    mq_close(q_server);
    mq_close(q_client);
    mq_unlink("CLIENT_ONE");
}
```