

# Fault Tolerance



ARCOS Group

Distributed Systems

Bachelor In Informatics Engineering

Universidad Carlos III de Madrid

# Introduction

The reliability of a system measures how well it conforms to its specification

A defect is a deviation from the expected external behavior

- Result of internal errors

The cause (mechanical, algorithmical, etc) of an error is a failure (fault)

- May be Hw or SW

# Examples of failures

## Ariane 5 explosion in 1996

- Sent by ESA in june of 1996 (first trip)
- Development cost: 10 yrs, \$7000 mil.
- Exploded 40s after takeoff, at 3700m
- The failure was due to the lost of height information
- Cause: software design error
- Conversion of real floating point nr on 64 bits to int on 16 bits results in value  $> 32767$  (max int on 16 bits) therefore exception.

# Examples of failures

## Patriot missiles

- Used in the 1991 Gulf war to intercept Scud missiles
- Failure due to error in timing
- Internal clk expresses tenths of sec as int, which gets converted to real on 24 bits (with corresponding loss in precision)

# Examples of failures

Viking probe sent to Venus

instead of Fortran `DO 20 I = 1,100`

(i.e. 100 iteration loop over label 20) it said

`DO 20 I = 1.100`

Compiler does not consider spaces so interprets as

`DO20I = 1.100`

(i.e. var declaration (`O20I`) with value 1.100)

D identifies a real type

# Examples of failures

## The Spirit robot

- Had a RAD6000 similar to an IBM RS6000 (with radiation protection) with an early-generation PowerPC processor w. 128 MB RAM and 256 MB flash memory
- Uses RTOS VxWorks within 32 MB RAM
- The very file needs some OS-controlled information
- The 32 MB ended up full of data associated to thousands of files and OS rebooted
- Every reboot ended up in memory full and subsequent reboot therefore systems seemed not working

# The moth inside the computer... bug

9/9

0800 Andam started  
1000 " stopped - andam ✓

		{ 1.2700	9.037847025
			9.037846995 convd
13 <sup>00</sup> (032)	MP - MC	<del>1.982647000</del> 2.130476415	<del>4.615925059 (-2)</del>
(033)	PRO 2	2.130476415	
	convd	2.130676415	

Relays 6-2 in 033 failed speed test  
in relay " " 11,000 test.

1100 Started Cosine Tape (Sine check)  
1525 Started Multi-Adder Test.

1545

Relay #70 Panel F  
(moth) in relay.

First actual case of bug being found.  
1630 Antangant started.  
1700 closed down.

1700 and tangent started.  
1700 closed down.

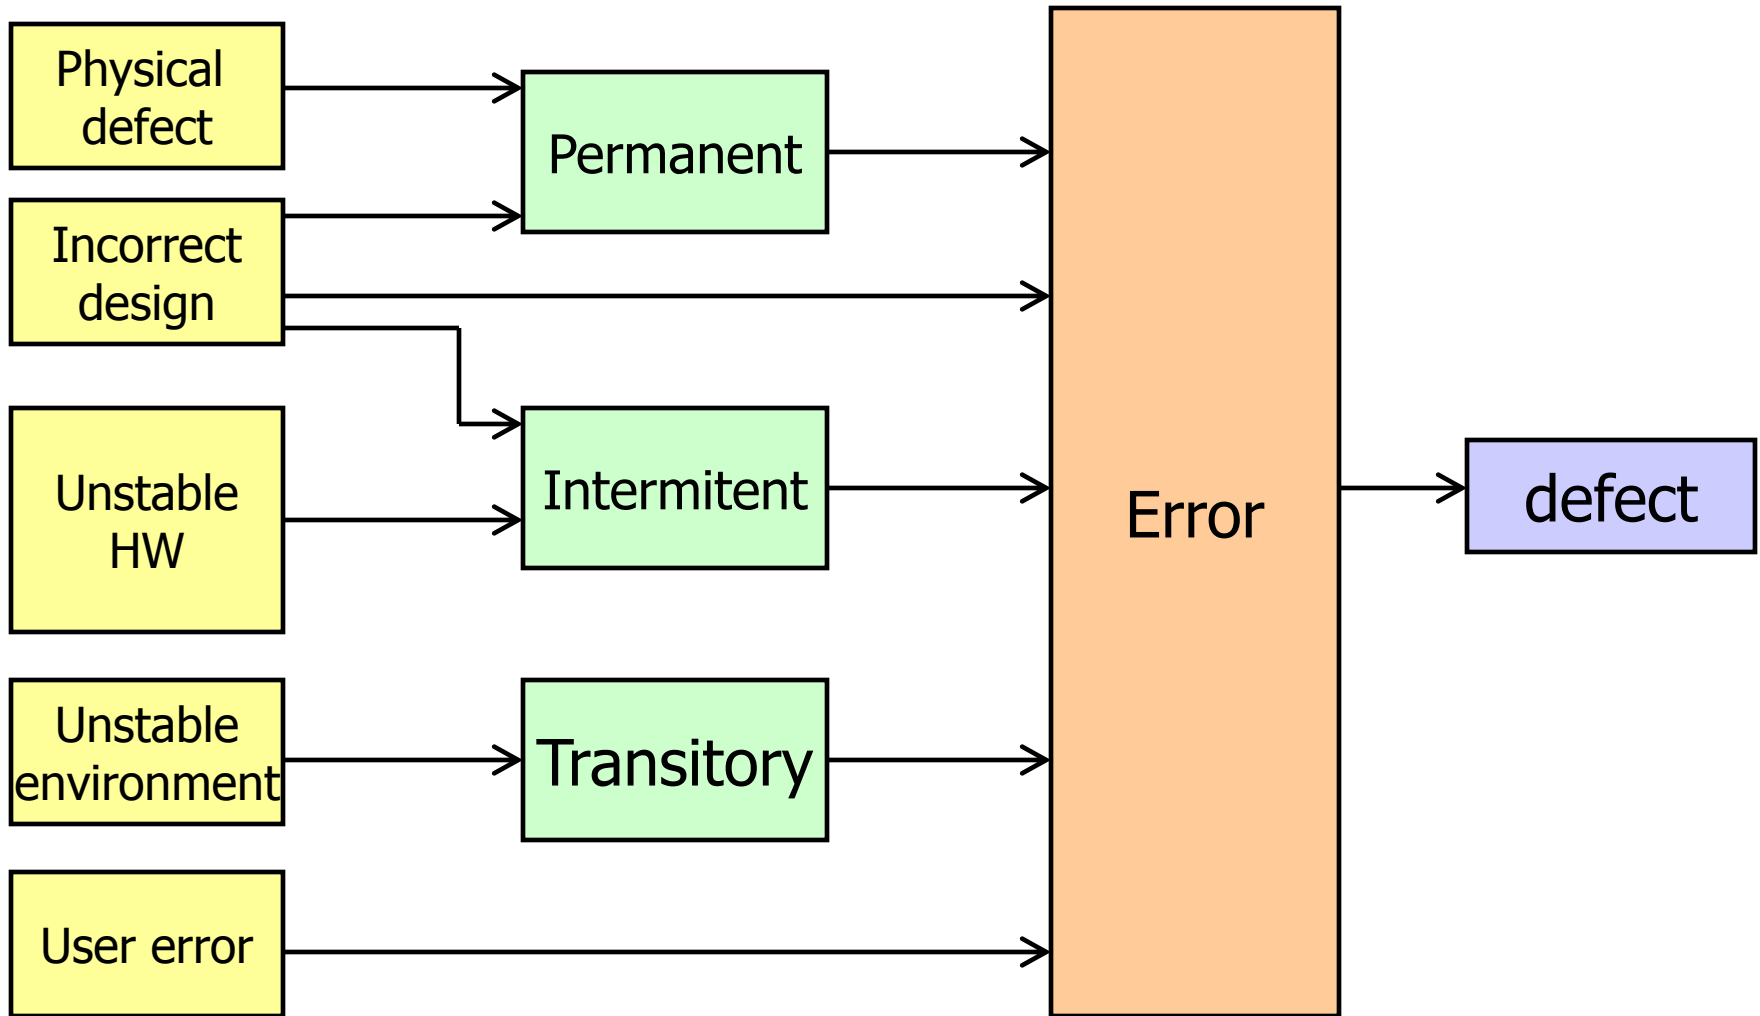
# Examples of failures

- More on SW failures at:

<http://www.cs.tau.ac.il/~nachumd/verify/horror.html>



# Failure types



# Definition

Fault tolerant system:

- A system capable to ensure the correct and uninterrupted execution despite HW / SW failures.

Goal:

- High reliability

Fault tolerance in distributed systems:

- Basic approach: Replication

# Replication

## Goals

- Better performance (caché)
- Better disponibility
  - If  $p$  is the failure prob. in a server
  - With  $n$  servers the failure prob. is  $p^n$

## Replication types

- Data
- Processes

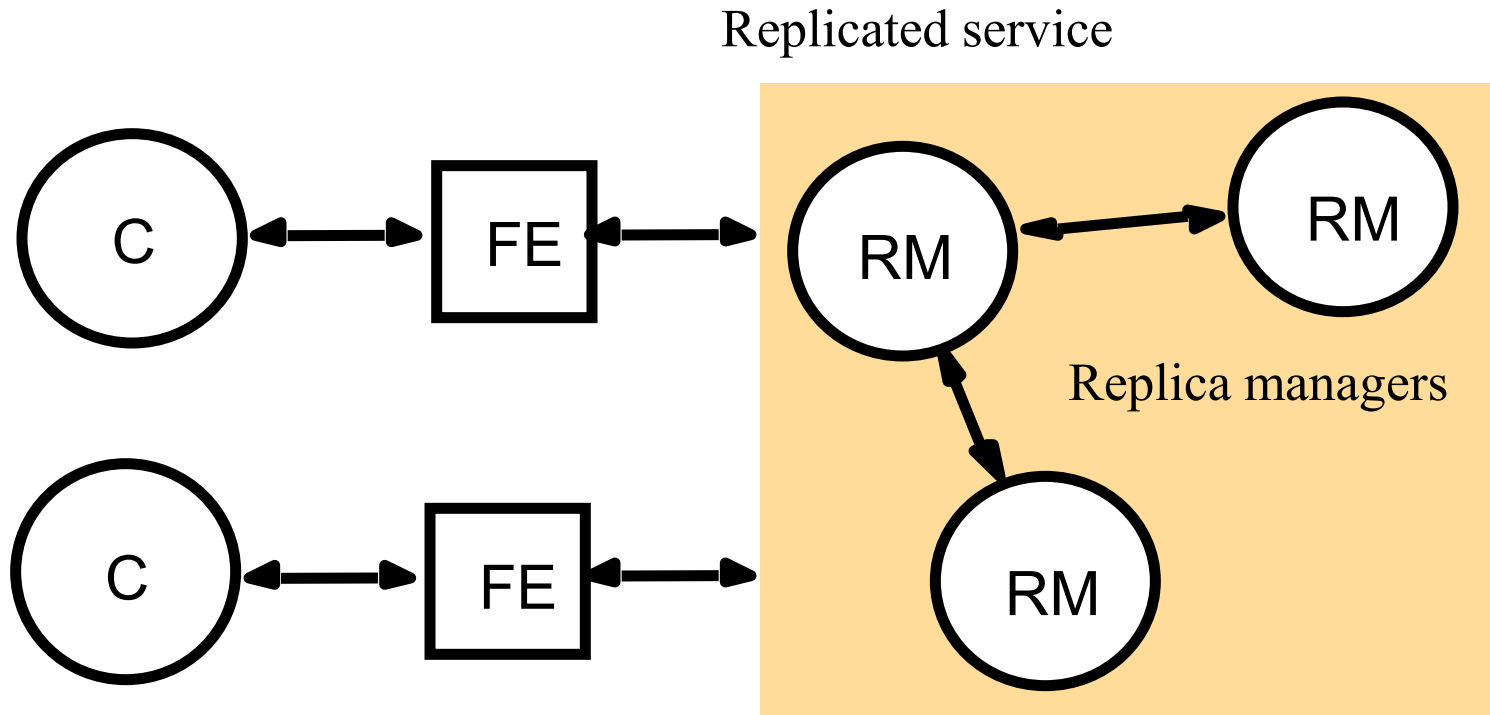
## Problems

- Consistency

## Requirements

- Transparency
- Consistency
- Performance

# Basic replication scheme



Front-end: makes replication transparent

# Replication methods

**Pesimistic:** ensure consistency of replicas

- Primary copy
- Active replicas
- Voting schemes (quorum)
  - Static
  - Dynamic

**Optimistic:** no consistency

- No limitations on data use when failure
- E.g.: version vectors, CODA file system

# Consistency models

## Strong:

- Pesimistic replication
- Total consistency

## Weak:

- Optimistic replication
- Local W w/o restrictions
- Modifications which result in inconsistencies must be rolled back or corrected
- Good for cases with few concurrent W processes

# Primary copy

If  $k$  failures one needs  $k+1$  copies

- Primary copy
- $K$  nodes w replicas



$R$ : send to primary

$W$ : send to primary

- Primary does the update
- And updates the rest of the replicas
- Then it replies to the client

When primary fails another node replaces it

# Replica synchronization






# Active replicas

All nodes act as primary

- Better performance on R

Atomic multicast necessary on W!

- W order is important



# Voting schemas (quorum)


Define two operations: **READ**, **WRITE**

**N** nodes that serve requests

- The **READ** operation is done on the **R** copies
- The **WRITE** operation is done on the **W** copies
- Every replica has a version number **V**
- It must be the case that:
  - **$R + W > N$**
  - **$W + W > N$**
  - **$R, W < N$**

# Examples of quorums

1



# How to choose W and R?

Two factors are important:

- Performance: depends on the % of R / W and their cost
  - $\text{Total cost} = \text{cost R} * p * \text{nr\_R} + \text{cost W} * (1 - p) * \text{nr\_W}$
- Fault tolerance: depends on the probability of failure
  - $\text{Prob. failure} = \text{Prob. failure R} + \text{Prob. failure W}$

E.g.:

- $N=7$
- Cost of W = 2 times cost of R (K)
- Percentage of R ( $p$ ) = 70%
- Prob. failure = 0.05

# Reliability for different configurations

Series configurations:

$R_i(t)$  = reliability of component  $i$

configuration fails whenever SOME component fails

assume independent failures

overall reliability:  $R(t) = \prod_{i=1}^N R_i(t)$




Parallel configurations:

ALL components must fail

overall reliability: 1 – prob. that all fail

$$R(t) = 1 - \prod_{i=1}^N Q_i(t) \quad \text{donde} \quad Q_i(t) = 1 - R_i(t)$$



# Solution

R	W	Coste	Probabilidad de fallo en R	Probabilidad de fallo en W	Probabilidad de fallo
1	7	4,9	7,8125E-10	0,301662704	9,05E-02
2	6	5	1,04688E-07	0,044380542	1,33E-02
3	5	5,1	6,02734E-06	0,003757043	1,13E-03
4	4	5,2	0,000193578	0,000193578	1,94E-04

# Voting methods

## READ

- Read all  $R$  copies, maintain last version

## WRITE

- First READ to figure out version number
- Initiate a 2PC to update data and version number in  $W$  copies

# *two-phase commit*

## Two-phase-commit (2PC)

There exists a coordinator

Coordinator:

multicast: *ok to commit?*

wait for replies

all ok => *send(commit)*


else => *send(abort)*

Processes:

*ok to commit* => save changes, reply *ok*

*commit* => commit changes

*abort* => erase changes






# Hierarchical voting

W grows with the number of replicas

Solution: hierarchical voting

- E.g.: replica nr =  $5 \times 5 = 25$  (leaves)



# Hierarchical voting

quorum on levels

R1	W1	R2	W2	RT	WT
1	5	1	5	1	25
1	5	2	4	2	20
1	5	3	3	3	15
2	4	2	4	4	16
2	4	3	3	6	12
3	3	3	3	9	9


Which one we choose?

# Dynamic methods

Previous methods (**static**) are not adaptative when failures occur

E.g.:

- 4 replicas with  $R=2$  y  $W=3$
- Given following partition



- Cannot write

# Dynamic voting

Each data  $d$  exists on  $N$  replicas  $\{d1..dn\}$

Each data  $d_i$  on node  $i$  has a version number  $VNi$   
(initially 0)

Actual version:  $NVA(d) = \max \{VNi\} \quad \forall i$

A replica of  $d_i$  is actual if  $VNi = NVA$

A group is a **majority partition** if it contains a  
majority of actual copies of  $d$

Each  $d_i$  has a cardinality update number  
associated  $SCi = \text{nr of nodes that participated}$   
in the update

# Dynamic voting

Initially  $SC_i = N$

When updating  $d_i$

- $SC_i =$  nr of copies of  $d$  updated during the current  $W$

A node may  $W$  if it is a member of a majority partition

# W algorithm

$\forall i$  requests  $NVi$  and  $SCi$

$M = \max \{NVi\}$  including itself

$I = \{i \text{ s.t. } NVi = M\}$

$N = \max \{SCi, i \in I\}$

**if  $|I| \leq N/2$**

then

node not member of majority partition, don't W

else {

$\forall \text{ nodes } \in I$

update

$VNi = M+1$

$SCi = \lceil I \rceil$

}

E.g.


N= 5

Initially:

	A	B	C	D	E
VN	9	9	9	9	9
SC	5	5	5	5	5

partition:

	A	B	C	D	E
VN	9	9	9	9	9
SC	5	5	5	5	5



# E.g.


W on partition 2?

- $M = \max\{9, 9\} = 9$
- $I = \{D, E\}$
- $N = 5, |I| = 2 \leq 5/2 \Rightarrow \text{nope}$

W on partition 1?

- $M = \max\{9, 9, 9\} = 9$
- $I = \{A, B, C\}$
- $N = 5$
- $|I| = 3 > 5/2 \Rightarrow \text{yep}$

	A	B	C	D	E
VN	10	10	10	9	9
SC	3	3	3	5	5





E.g.

New partition

	A	B	C	D	E
VN	10	10	10	9	9
SC	3	3	3	5	5

Partition 1      Partition 2      Partition 3

W on partition1?

- $N = \max\{10, 10\} = 10$
- $I = \{A, B\}$
- $N = 3$
- $|I| = 2 > 3/2 \Rightarrow \text{yep}$

E.g.

	A	B	C	D	E
VN	11	11	10	9	9
SC	2	2	3	5	5

Partition 1      Partition 2      Partition 3

# When a node joins a group

...it must update its state:

$$M = \max \{VN_i\}$$

$$I = \{A_j, \text{ s.t. } M = VN_j\}$$

$$N = \max \{SC_k, k \in I\}$$

**if  $|I| \leq N/2$**

then

cannot join

else {

update state

$$VN_i = M$$

$$SC_i = N + 1$$

}

E.g.

join 2 and 3

	A	B	C	D	E
VN	11	11	10	9	9
SC	2	2	3	5	5

Partition 1      Partition 2

W to partition 2?

- $M = \max\{10, 9, 9\} = 10$
- $I = \{C\}$
- $N = 3$
- $|I| = 1 \leq 3/2 \Rightarrow$  nope

E.g.

Join 1 and 2

	A	B	C	D	E
VN	11	11	10	9	9
SC	2	2	3	5	5

Partition 1

Partition 2

W to partition 1?

- $M = \max\{11, 11, 10\} = 11$
- $I = \{A, B\}$
- $N = 2$
- $|I| = 2 > 2/2 \Rightarrow \text{yep}$

E.g.

	A	B	C	D	E
VN	12	12	12	9	9
SC	3	3	3	5	5

Partition 1

Partition 2

# Replication of the CODA file system

## Optimistic replication

Every copy has a version vector  $V$  w.  $n$  components = replication degree

For node  $i$   $V_i[j]$  is the number of updates on  $j$ 's copy

W/o network failures all vectors are the same, otherwise they differ

Given  $V1$  &  $V2$ ,  $V1$  dominates  $V2$  iff  $V1(i) \geq V2(i) \forall i$

If  $V1$  dominates  $V2$  there are more updates in the copy of  $V1$

$V1$  &  $V2$  in conflict if neither dominates

When groups join they compare vectors

- The copy of group w the dominant vector updates the second
- If conflicts then error

E.g.

3 servers  $\{A, B, C\}$

Initially  $V = (0,0,0)$  for all

On update:  $V=(1,1,1)$  for all

Assume network failure: Group 1:  $\{A,B\}$ , Group 2:  $\{C\}$

Update group 1:  $V=(2,2,1)$  for group 1

Assume network failure: Group 1:  $\{A\}$ ,  $V=(2,2,1)$

Group 2:  $\{B, C\}$

- $(2,2,1) \geq (1,1,1) \Rightarrow$  update C's copy,  $V = (2,2,2)$  for B, C

Update group 2:  $V=(2,3,3)$  for  $\{B,C\}$



## E.g. (cont'd)

Case 1: join  $\{A\}$  and  $\{B,C\}$

- $(2,2,1) \leq (2,3,3) \Rightarrow$  update  $\{A\}$ 's copy,  $V=(3,3,3)$

Case 2:

- Modify  $\{A\}$ 's version  $\Rightarrow$  en A,  $V=(3,2,1)$
- Join A w.  $V=(3,2,1)$  with  $\{B,C\}$  w.  $V=(2,3,3)$
- Compare  $(3,2,1)$  and  $(2,3,3)$  neither dominates!  
**conflict**