



Processes and thread management

Operating Systems Design


References

- Silberschatz **Operating System Concepts** *Ninth Edition* 2012. Chapters 3, 4, and 6

Process Concept

- An operating system executes a variety of programs:
 - Batch system – jobs
 - Time-shared systems – user programs or tasks
- Textbook uses the terms *job* and *process* almost interchangeably
- Process – a program in execution; process execution must progress in sequential fashion
- A process includes:
 - program counter
 - stack
 - data section

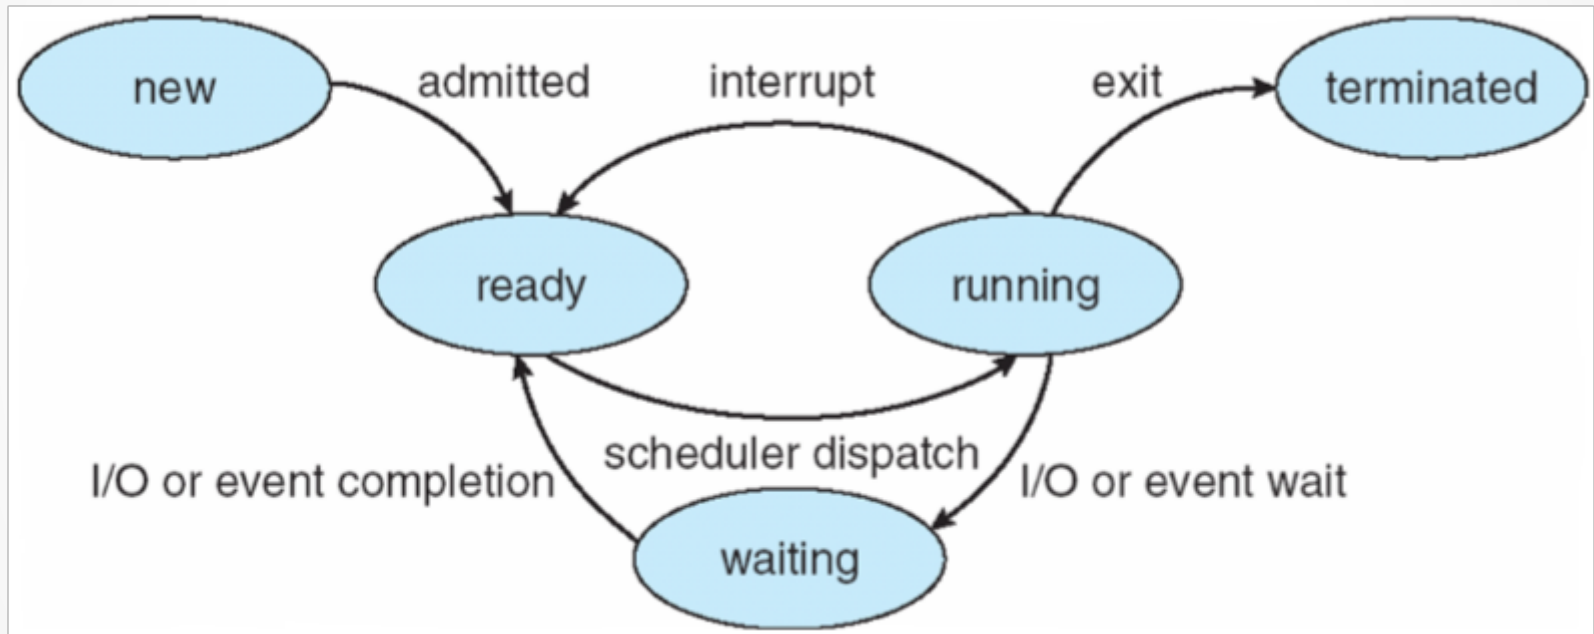
Process in Memory



Process State

- As a process executes, it changes *state*
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution

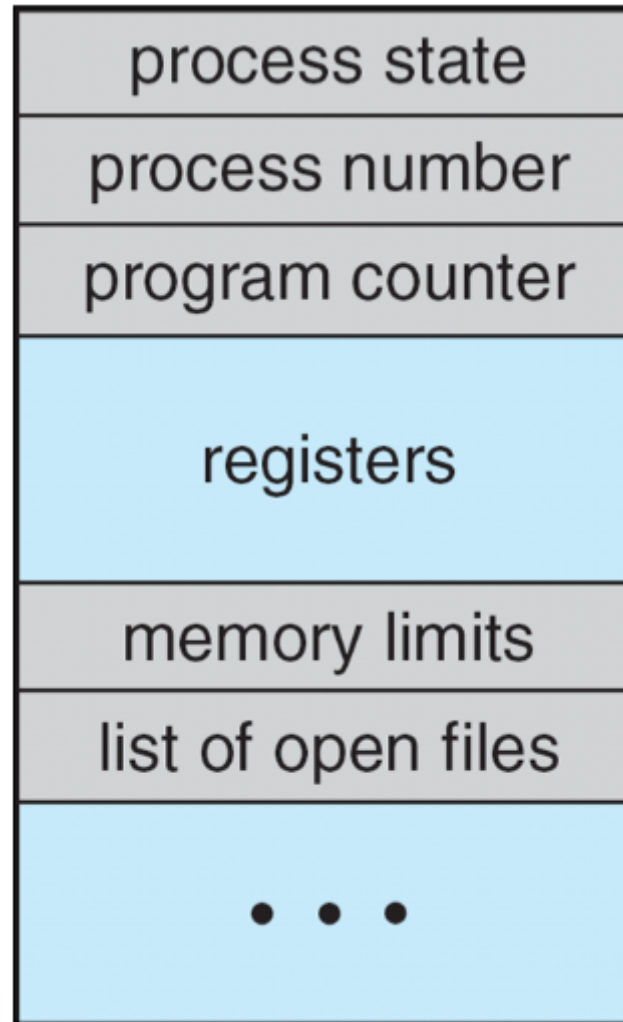
Diagram of Process State



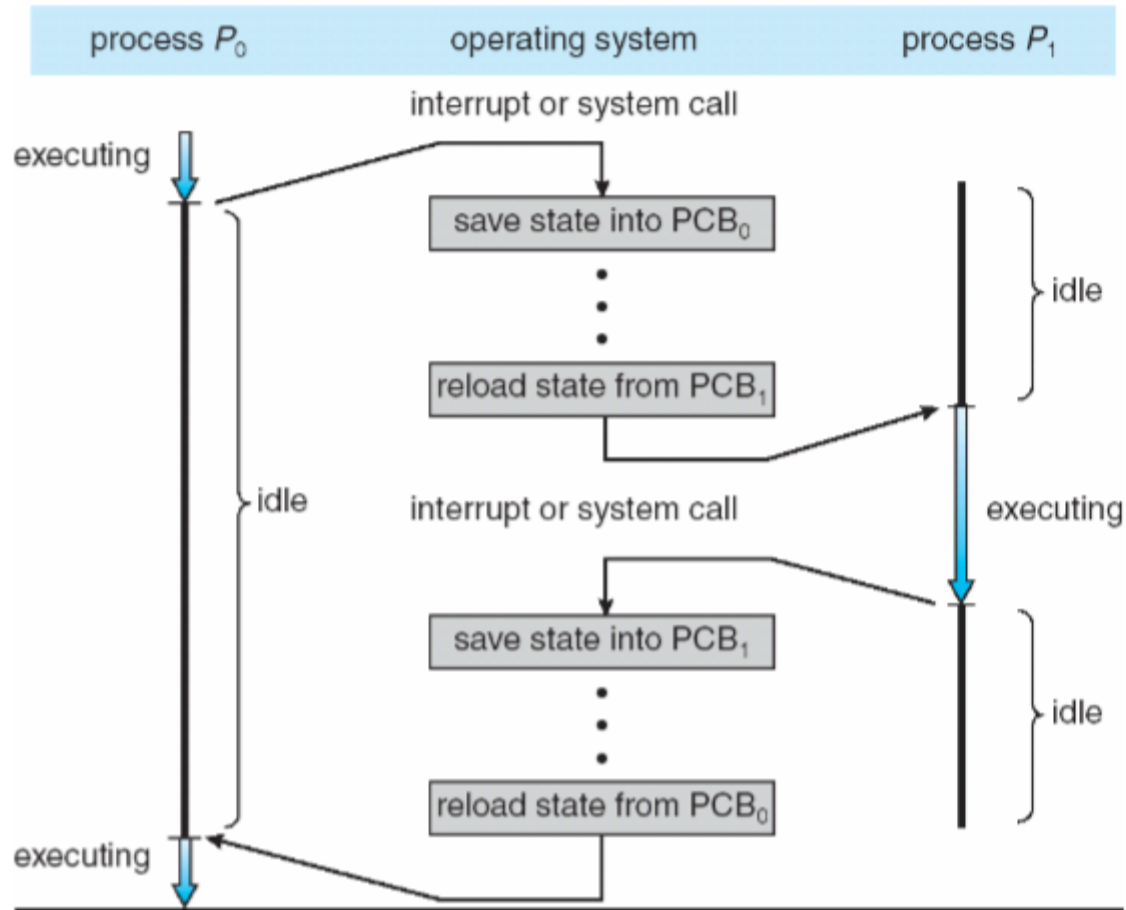
Process Control Block (PCB)

- Information associated with each process
 - Process state
 - Program counter
 - CPU registers
 - CPU scheduling information
 - Memory-management information
 - Accounting information
 - I/O status information
- PCBs are stored in a global **process table**

Process Control Block (PCB)



CPU Switch From Process to Process

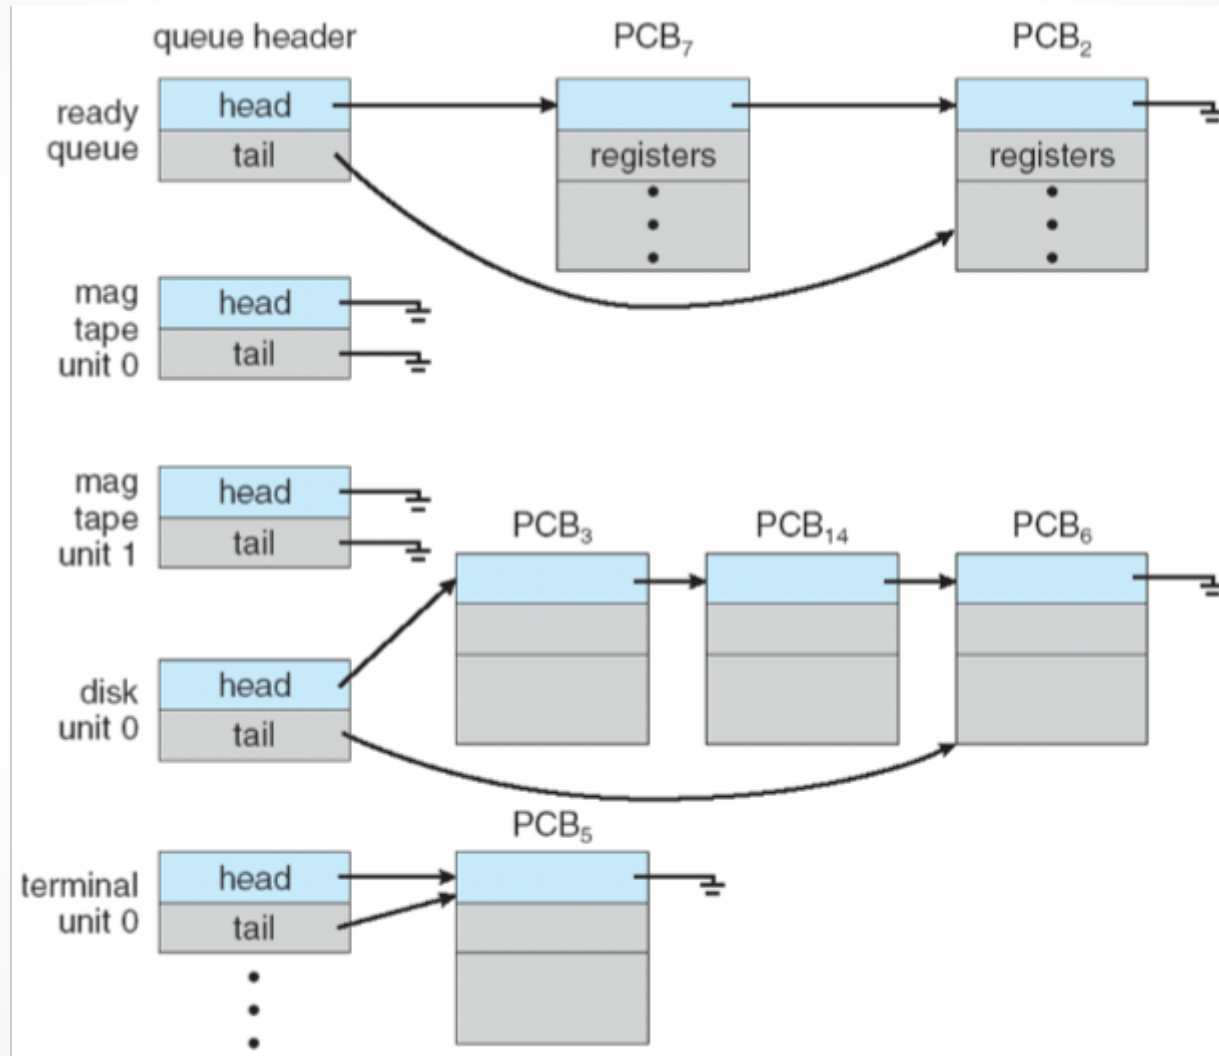


Process Scheduling Queues

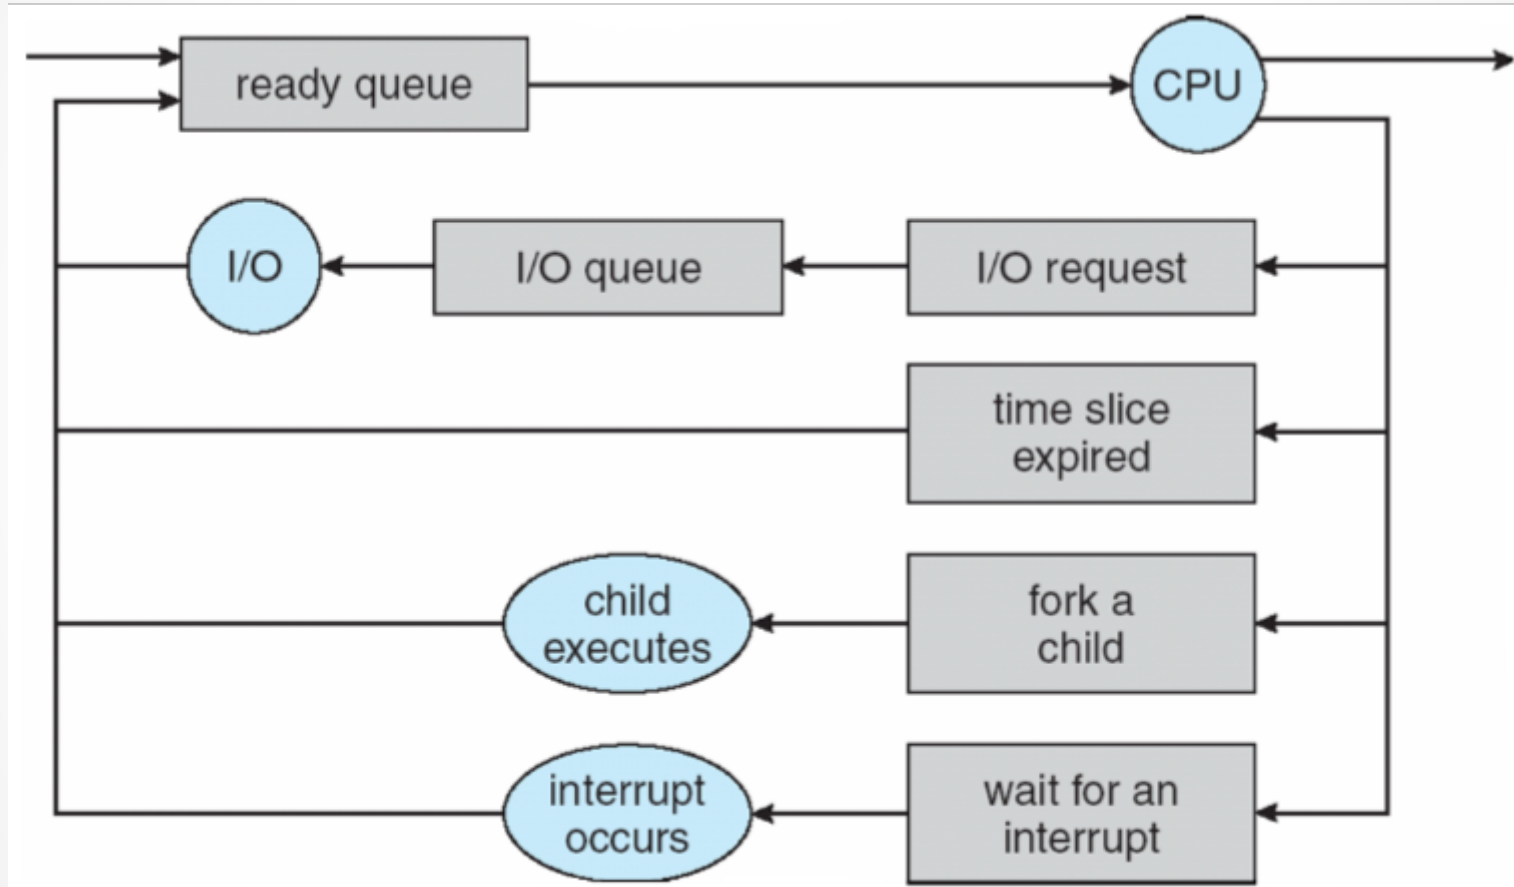
- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- Processes migrate among the various queues



Ready Queue And Various I/O Device Queues



Representation of Process Scheduling

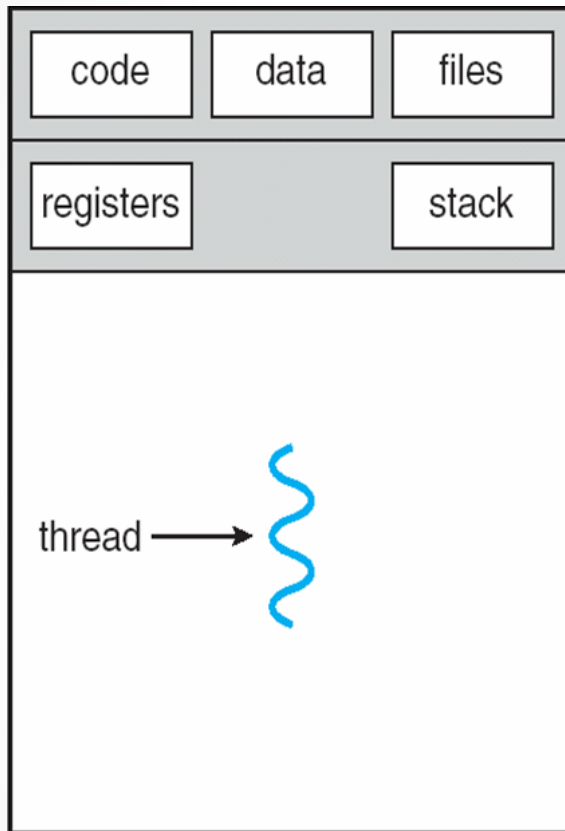


Threads and processes

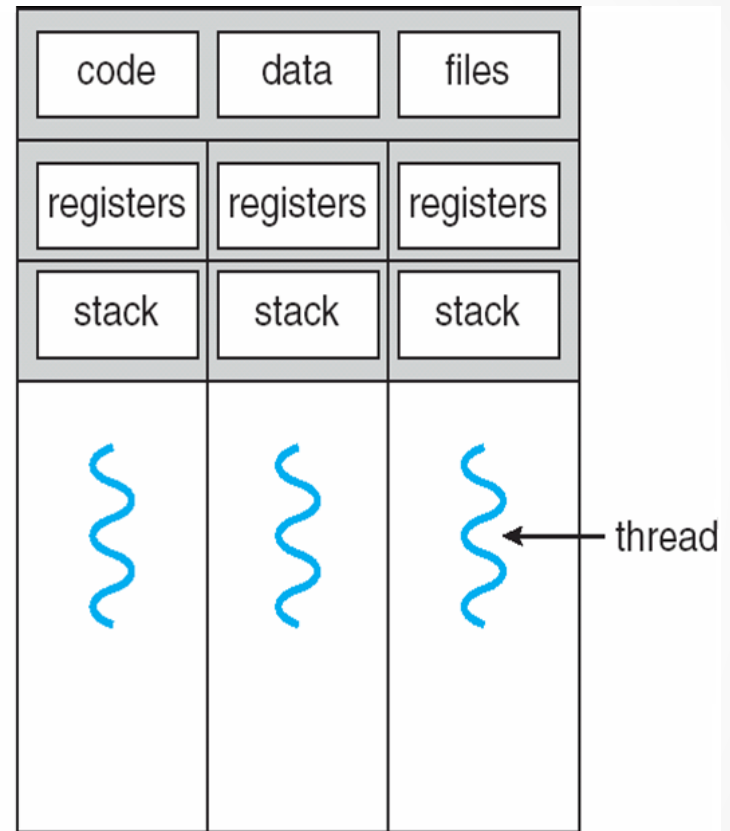
- Thread = a lightweight process
- Threads of one process share
 - Data section (heap)
 - Code
 - Open files
 - Signals
- Each thread of a process has his own
 - Stack
 - Register state



Single and Multithreaded Processes



single-threaded process



multithreaded process



Threads

- Two types of thread implementations
 - User threads
 - Kernel threads
- A **thread library** provides programmer with API for creating and managing threads
- In the first programming assignment you will implement a **user thread library**

User Threads

- Thread management done by user-level thread library
- No kernel support
 - The kernel is not aware of the existence of several threads
 - If one blocks for I/O the kernel may infer that the whole process blocks for I/O and schedule another process



Kernel Threads

- Supported by the Kernel
 - If one thread within one process blocks for I/O the kernel knows how to schedule another thread from the same process
- Most OS support kernel threads
 - Windows
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X



Scheduler

- Selects from among the processes/threads in memory that are ready to execute, and allocates the CPU to one of them
- CPU scheduling decisions may take place when a process/thread:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive (voluntary)**
- All other scheduling is **preemptive (involuntary)**



Dispatcher (activator)

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- Dispatch latency – time it takes for the dispatcher to stop one process and start another running



Scheduling Criteria

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)



Scheduling Algorithm Optimization

Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

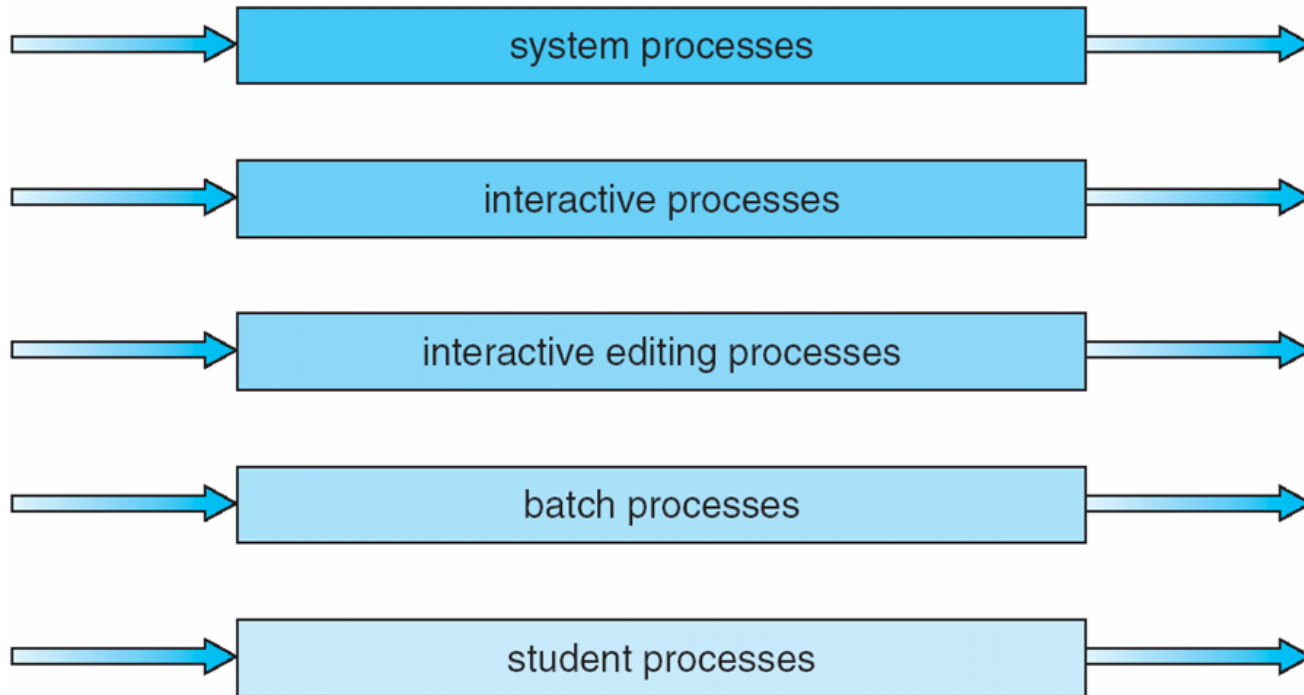


Main scheduling algorithms

- FCFS (non-preemptive)
- Round-Robin (preemptive)
- SJF
 - Preemptive
 - Non-preemptive
- Priority-based
 - Preemptive
 - Non-preemptive
- Multi-level queue

Multilevel Queue Scheduling

highest priority



lowest priority



Task context and scheduling

- The context of a user task is standardized by POSIX

```
typedef struct ucontext {  
    struct ucontext uc_link; /* context that will be resumed  
when the current context terminates*  
    sigset_t          uc_sigmask; /* set of signals blocked  
in this context */  
    stack_t           uc_stack; /* stack used by this  
context */  
    mcontext_t        uc_mcontext; /* machine-specific  
representation of  
the saved context */  
    ...  
} ucontext_t;
```


Manipulation of the user context

- `Getcontext(ucontext_t *ucp)`: initializes a context
- `Setcontext(const ucontext_t *ucp)`: starts to run the ucp context on the CPU
- `Makecontext(ucontext_t *ucp, void *func(), int argc, ...)`: creates a new thread using the ucp context previously created by `getcontext`
- `Swapcontext(ucontext_t *old_ucp, ucontext_t *new_ucp)`: starts to run the new_ucp context on the CPU and saves the currently running context into old_ucp